

# Blockchains & Distributed Ledgers

Lecture 02

Aggelos Kiayias

Slide credits: AK, Dionysis Zindros, Christos Nasikas

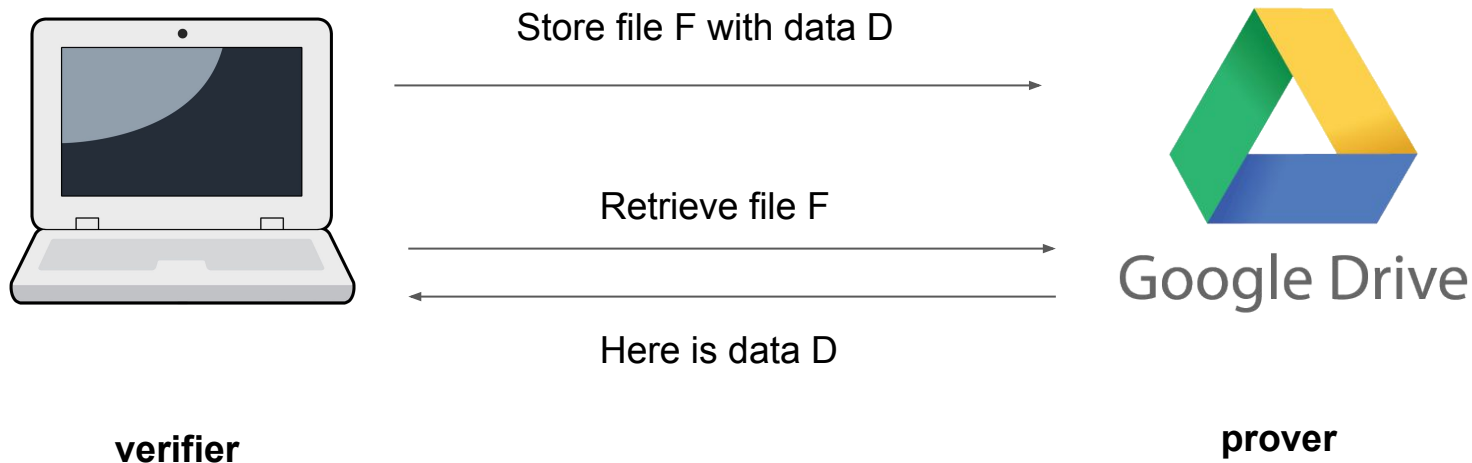
# Overview

- Motivation: Server file storage
- Merkle trees to store lists
- Proofs-of-inclusion
- Merkle trees to store sets
- Proofs-of-non-inclusion
- Merkle–Patricia tries to store key:value pairs
- Blocks and blockchains

# Authenticated Data Structures

- Like regular data structures, but cryptographically authenticated
- Allows a **verifier** to store, retrieve and operate on data with an untrusted **prover**

# The file storage problem



# The file storage problem

- Client wants to store a file on a server
- File has a name  $F$  and data  $D$
- Clients wants to retrieve file  $F$  later

# File storage: Basic protocol

- Client sends file  $F$  with data  $D$  to server
- Server stores  $(F, D)$
- Client deletes  $D$
- Client requests  $F$  from server
- Server returns  $D$
- Client has recovered  $D$

# File storage: Protocol against adversaries

- What if **server is adversarial** and returns  $D' \neq D$ ?

# File storage: Protocol against adversaries

Trivial solution:

- Client does not delete  $D$
- When server returns  $D'$ , client compares  $D$  and  $D'$

...what if client doesn't have enough memory to store  $D$  for a long time?



# File storage: Hash-based protocol


- Client sends file F with data D to server
- Server stores (F, D)
- Client stores  $H(D)$ , deletes D
- Client requests F from server
- Server returns  $D'$
- Client compares  $H(D') = H(D)$

# File storage: File chunks

- What if client wants to retrieve the 200,019th byte of the file?
- Must download the whole file...
- Merkle trees to the rescue!


# Merkle Tree

- An **authenticated** binary tree
- Split file into **chunks** of, say, 1KB




# Merkle Tree

- **Hash** each chunk using a cryptographic hash function (SHA256)
- Convention: Arrows show direction of hash function application




# Merkle Tree

- **Combine** them by two to create a binary tree
- Each node stores the **hash** of the **concat** of its children



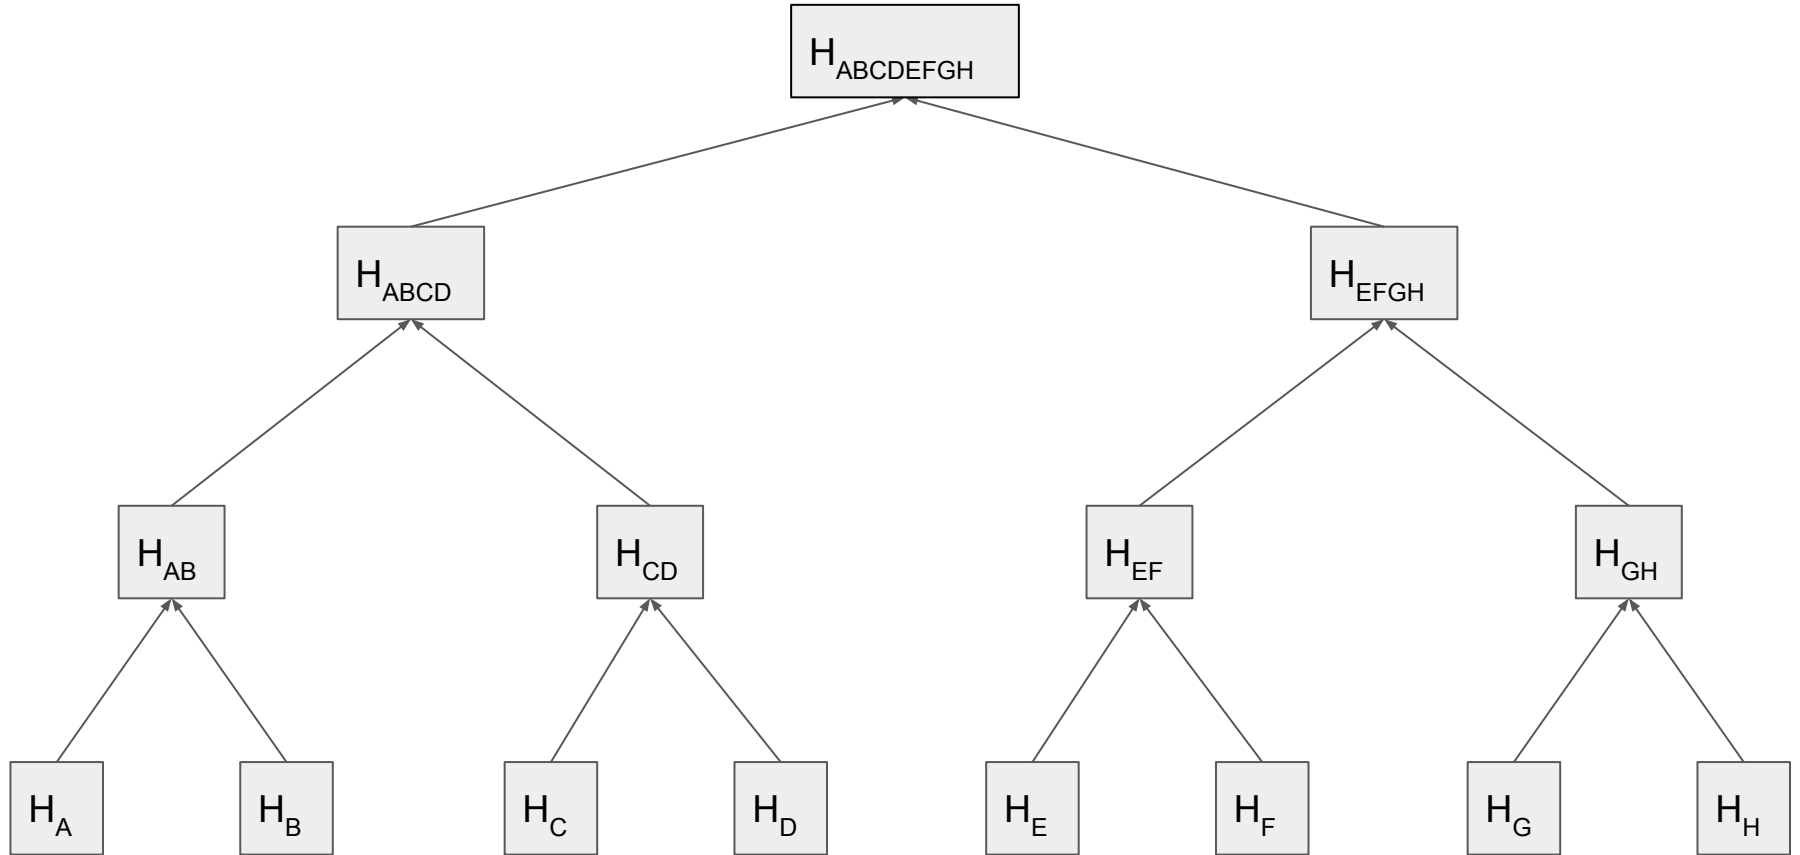
# Merkle Tree



# Proofs-of-inclusion

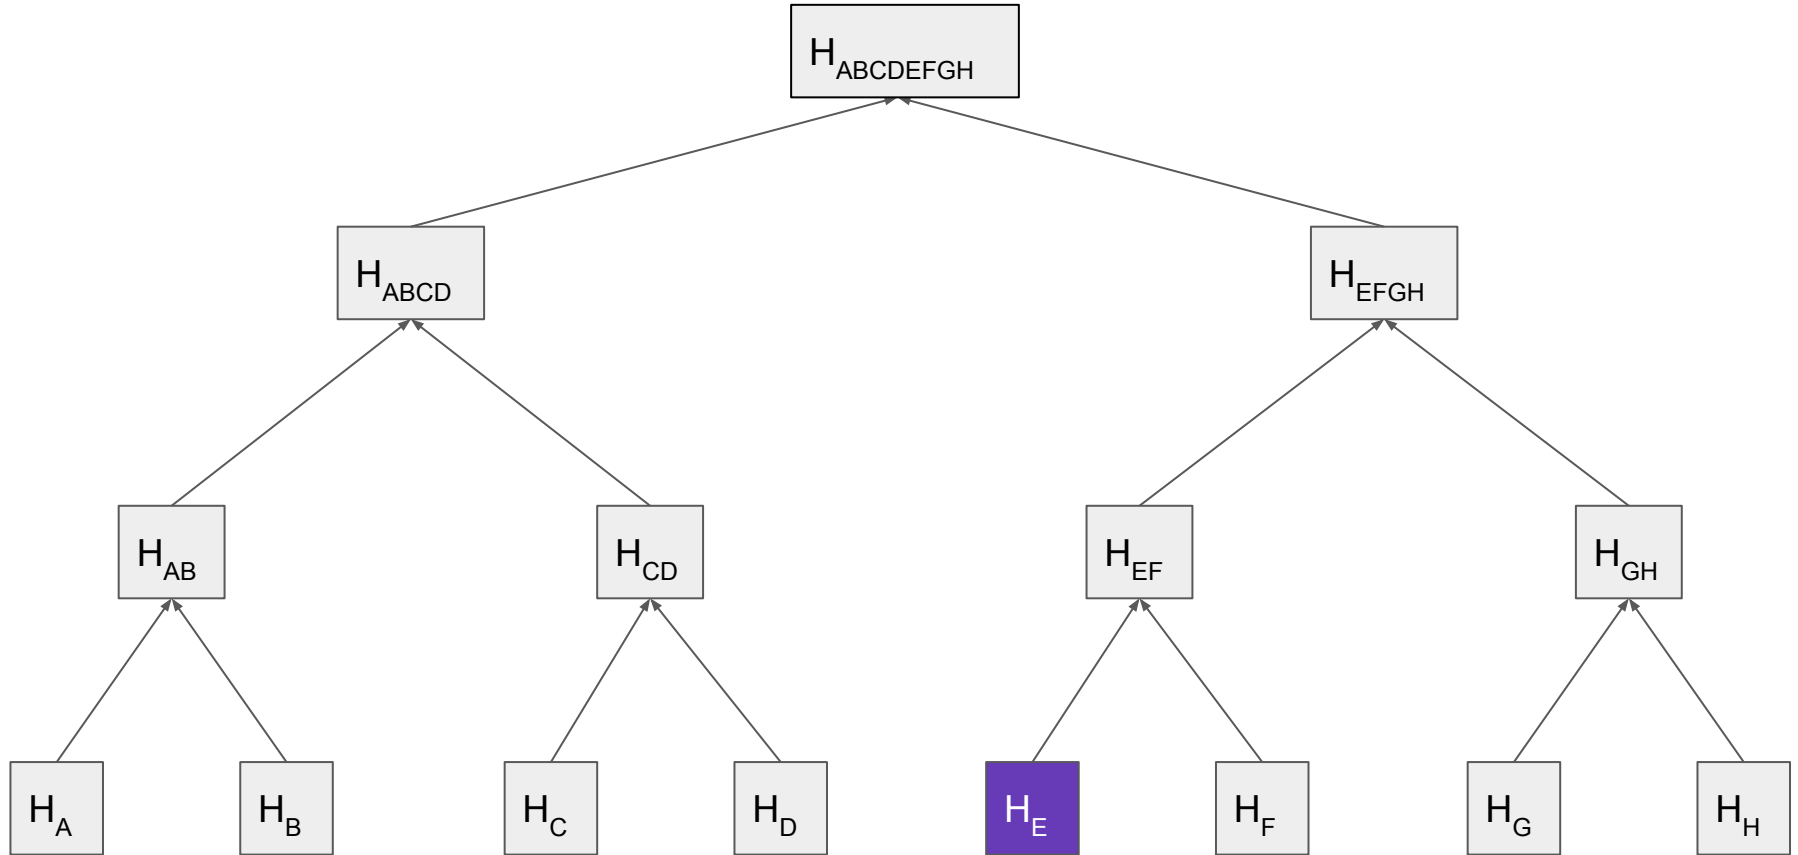
- Client creates Merkle Tree root **MTR** from initial file data  $D$
- Client sends file data  $D$  to server
- Client deletes data  $D$ , but stores MTR (32 bytes)
- Client requests chunk  $x$  from server
- Server returns chunk  $x$  and short proof-of-inclusion  $\pi$
- Client checks that chunk  $x$  is included in MTR using proof  $\pi$

# Merkle tree: proof of inclusion






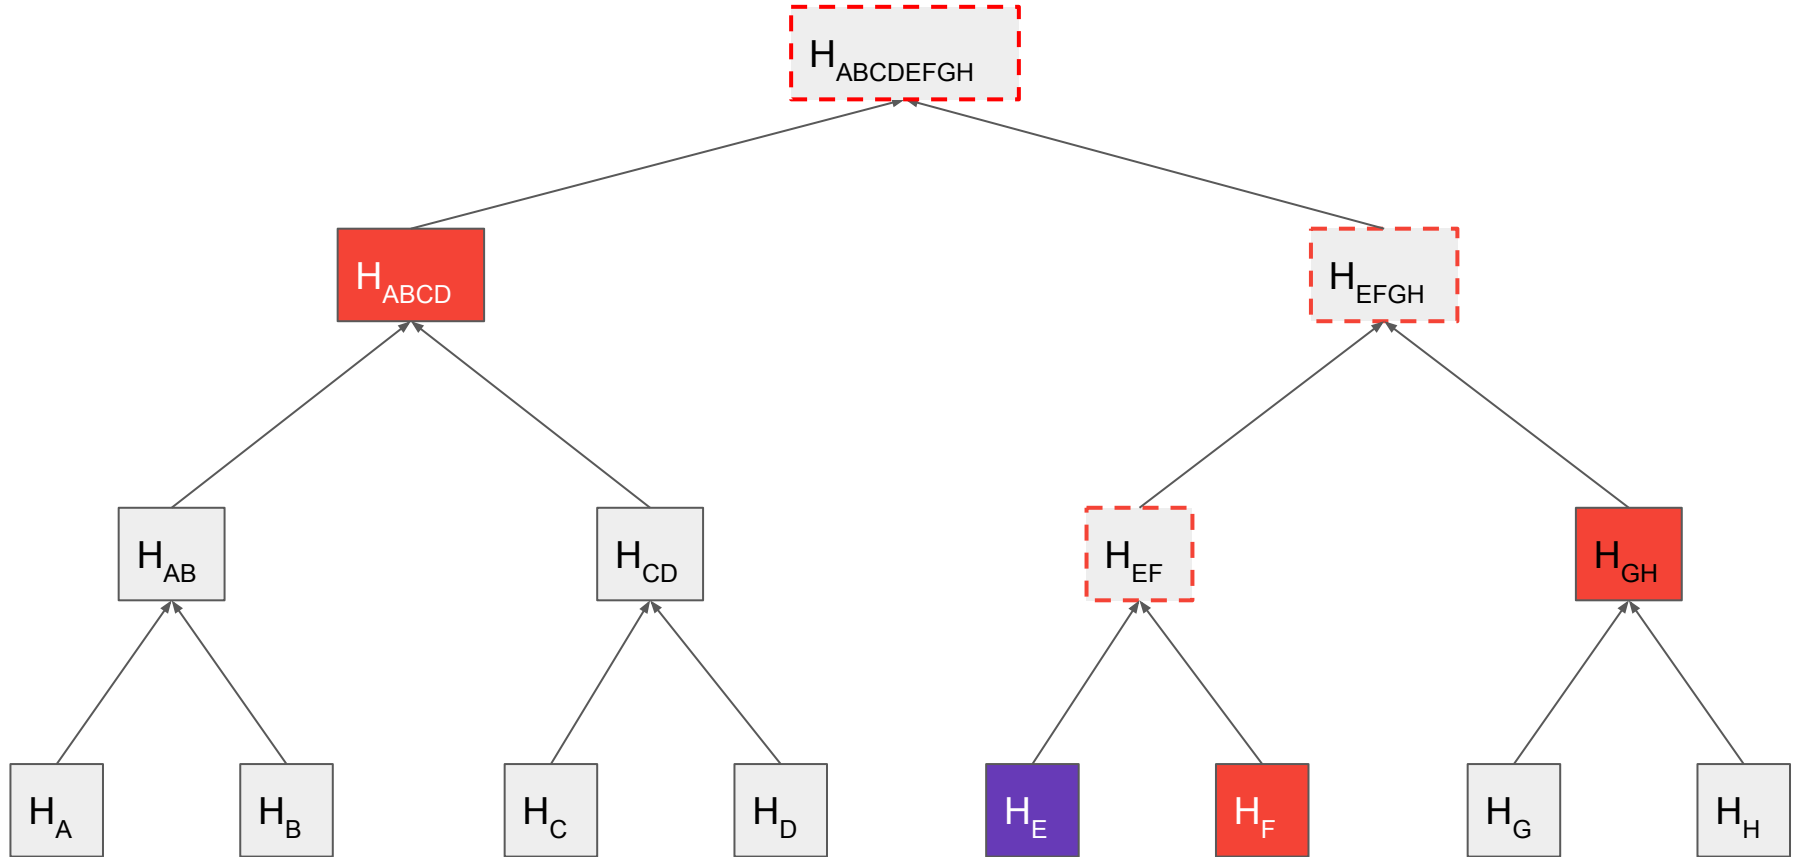
# Merkle tree: proof of inclusion



# Merkle tree: proof of inclusion



# Merkle tree: proof of inclusion



# Merkle Tree proof-of-inclusion

- Prover sends chunk
- Prover sends **siblings** along path connecting leaf to MTR
- Verifier computes hashes along the path connecting leaf to MTR
- Verifier checks that computed root = MTR
- How big is proof-of-inclusion?

# Proof-of-inclusion succinctness

$$|\pi| \in \Theta(\lg|D|)$$

# Merkle Tree proof-of-inclusion security

- If adversary can present proof-of-inclusion for incorrect leaf, then we can break the hash function
- Proof is by computational reduction (whiteboard)

# Merkle Tree protocol

MT-construct(D)

- Construct a Merkle Tree with given data D
- Returns the Merkle Tree root
- If  $|D| = \text{chunk size}$ , then:  $\text{MT-construct}(D) = H(D)$
- Otherwise:

$$\text{MT-construct}(D) = H(\text{MT-construct}(D_1) \parallel \text{MT-construct}(D_2))$$

$$\text{where } D = D_1 \parallel D_2$$

# Merkle Tree protocol

MT-prove( $D$ ,  $x$ )

- Given data  $D$  and element  $x$  in  $D$ , construct proof-of-inclusion
- Returns the proof-of-inclusion  $\pi$  to be used with MT-construct( $D$ )
- Proof contains:
  - Siblings on path connecting  $x$  to root
  - A bit for each sibling indicating whether the path we are taking is left or right



# Merkle Tree protocol

$\text{MT-verify}(r, \pi, x)$

- Given Merkle Tree root  $r$ , element  $x$ , and proof-of-inclusion  $\pi$
- Outputs true/false based on whether verification was successful

## Correctness

For all  $D, x$ :

$\text{MT-verify}(\text{MT-construct}(D), \text{MT-prove}(D, x), x) = \text{True}$

(Proof by direct application of hashes on path)

# Proof-of-inclusion security

- Assume the hash function is **collision-resistant**
- Collision resistance formal definition:
  - $\forall$  PPT  $A$ :  $\exists \text{negl}$ :  $\Pr[\text{coll-find}_{A,H}(\lambda)] \leq \text{negl}(\lambda)$
- Where coll-find is the collision finding game:

```
def coll-findA,H(λ):  
    x1, x2 ← A(1λ)  
    if x1 ≠ x2 ∧ H(x1) = H(x2):  
        return 1  
    return 0
```

# Threat modelling with bad events

- When defining a security property precisely, specify what **bad event** we are trying to avoid
- In this case, the construction of a proof about a non-existent element
- It is important to allow the adversary to **choose** which Merkle Tree to attack
- It is possible that the vast majority of trees are not attackable...
- Hence, we define a **game** where the adversary chooses a data set  $D$  to construct the tree from, an element  $x$ , and a proof of  $\pi$
- The adversary can construct these arbitrarily.  $\pi$  does not need to be produced out of a tree!

# The Merkle Tree forgery game

```
def MT-forgeryA, Π(H)(λ):  
    (D, x, π) ← A(1λ)  
    if MT-verify(MT-construct(D), π, x) ∧ x ∉ D:  
        return 1  
    return 0
```

# The Merkle Tree security

$$\forall \text{ PPT } A: \exists \text{ negl}: \Pr[\text{MT-forgery}_{A, \Pi(H)}(\lambda)] \leq \text{negl}(\lambda)$$

# The theorem: Assumption $\rightarrow$ Desirable

**Theorem:** If  $H$  is collision-resistant, then the MT constructed from  $H$  is secure:

$$\forall \text{ PPT } A: \exists \text{ negl}: \Pr[\text{coll-find}_{A,H}(\lambda)] \leq \text{negl}(\lambda)$$

$\rightarrow$

$$\forall \text{ PPT } A: \exists \text{ negl}: \Pr[\text{MT-forgery}_{A,\Pi(H)}(\lambda)] \leq \text{negl}(\lambda)$$

# Proof strategy: Contraposition

By reductio ad absurdum using contraposition:

- Suppose for contradiction that  
**not**  $\forall$  PPT  $A$ :  $\exists$   $\text{negl}$ :  $\Pr[\text{MT-forgery}_{A, \Pi(H)}(\lambda)] \leq \text{negl}(\lambda)$   
i.e., the Merkle Tree construction is **not** secure
- It suffices to show that  
**not**  $\forall$  PPT  $A$ :  $\exists$   $\text{negl}$ :  $\Pr[\text{coll-find}_{A, H}(\lambda)] \leq \text{negl}(\lambda)$   
i.e., the hash function is **not** collision-resistant

# Proof strategy: Contraposition

- Suppose for contradiction that  
 $\exists \text{ PPT } A: \Pr[\text{MT-forgery}_{A, \Pi(H)}(\lambda)] \text{ is non-negl}$
- It suffices to show that  
 $\exists \text{ PPT } A^*: \Pr[\text{coll-find}_{A, H}(\lambda)] \text{ is non-negl}$


The PPT  $A$  is arbitrary, so we must use it as black box.

We show the existence of  $A^*$  by construction.

Since  $A$  is a machine, we can have  $A^*$  call  $A$  in its code.




# Proof strategy: Computational reduction

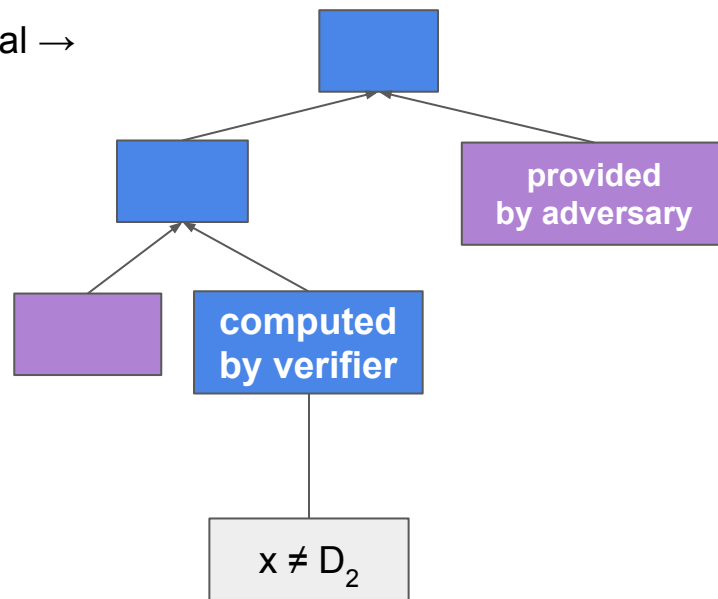


# Situation if adversary A wins


Real Merkle tree of D




Proof  $\pi$  provided by A



falsely claimed element



← roots are equal →




← leafs not equal →

Induction:


Take **lowest** tree level where nodes are **equal**

That level must exist, as roots are equal

That level cannot be a leaf, as leafs are not equal





← leafs not equal →



$$H(L^a \parallel R^a) = H(L^b \parallel R^b) \text{ but } R^a \neq R^b$$

We can extract a hash collision!



# Proof conclusion

$A^*$  works as follows:

- Checks if  $A$  has found forgery
- If not, aborts
- If yes, finds minimum level where hashes are equal
- This gives a hash collision

If  $A$  finds a MT forgery, then  $A^*$  finds a hash collision

equal by computational reduction

$$\underbrace{\Pr[\text{MT-forgery}_{A, \Pi(H)}(\lambda)]}_{\text{non-negligible by for-contradiction assumption}} = \underbrace{\Pr[\text{coll-find}_{A^*, H}(\lambda)]}_{\text{non-negligible, therefore contradiction}}$$

non-negligible by for-contradiction assumption

non-negligible, therefore contradiction

# Proof of security is nuanced

**Don't roll your own crypto.**

**Use standard code by others which you know is secure.**

# Merkle tree applications

- Bitcoin uses Merkle trees to store transaction
- BitTorrent uses Merkle Tree to exchange files
- Ethereum uses Merkle–Patricia tries for storage and transactions



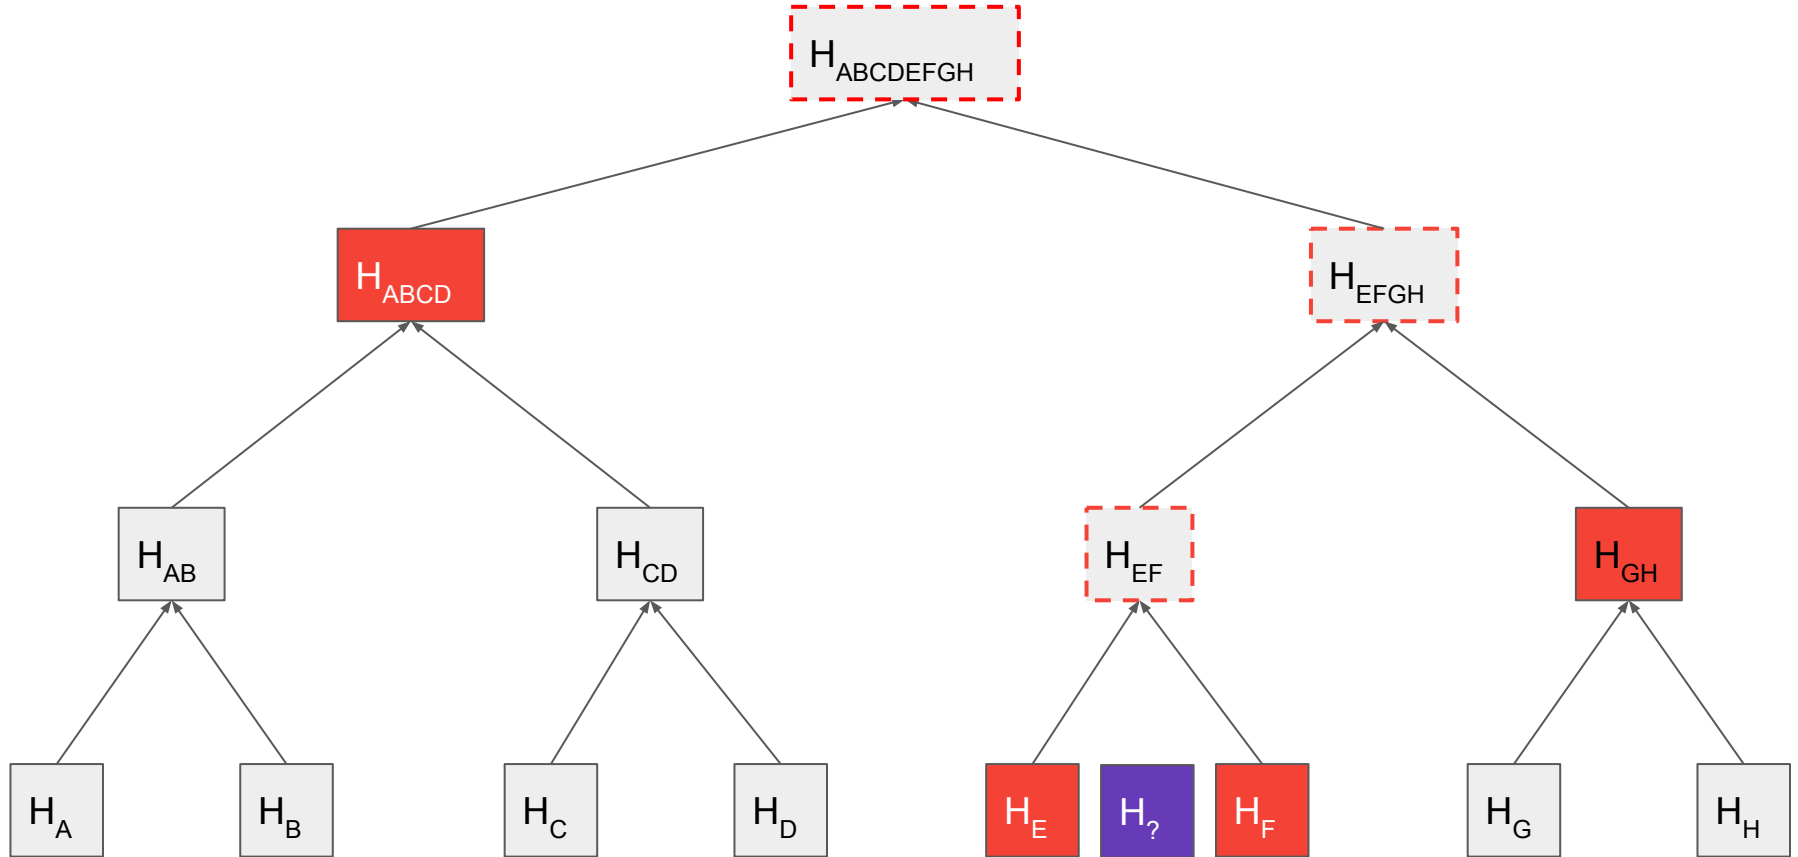
# Storing *sets* instead of lists

- Merkle Trees can be used to store sets of keys instead of lists
- Verifier asks prover to store a set of keys
- Verifier deletes set
- Verifier later asks prover if key belongs to set
- Prover provides proof-of-inclusion or proof-of-non-inclusion
- Prover can be adversarial

# Merkle trees for set storage

- Verifier sorts set elements
- Creates MTR on sorted set
- Proof-of-inclusion as before
- Proof-of-non-inclusion for  $x$ 
  - Show proof-of-inclusion for previous  $H_{<}$  and next  $H_{>}$  element in set
- Verifier checks that  $H_{<}$ ,  $H_{>}$  proofs-of-inclusion are correct
- Verifier checks that  $H_{<}$ ,  $H_{>}$  are adjacent in tree
- Verifier checks that  $H_{<} < x$  and  $H_{>} > x$
- The two proofs-of-inclusion can be compressed into one

# Merkle tree: proof of inclusion



# Tries

- Called also radix tree or prefix tree
- Search tree: ordered tree data structure
- Used to store a set or an associative array (key/value store)
- Keys usually are strings

# Tries

- Supports two operations: **add** and **query**
- **add** adds a string to the set
- **query** checks if a string is in the set (true/false)
- **Initialize**: Start with empty root

# Tries: add(string)

- Start at root
- Split string into characters
- For every character, follow an edge labelled by that character
- If edge does not exist, create it
- Mark the node you arrive at

# Tries: query(string)

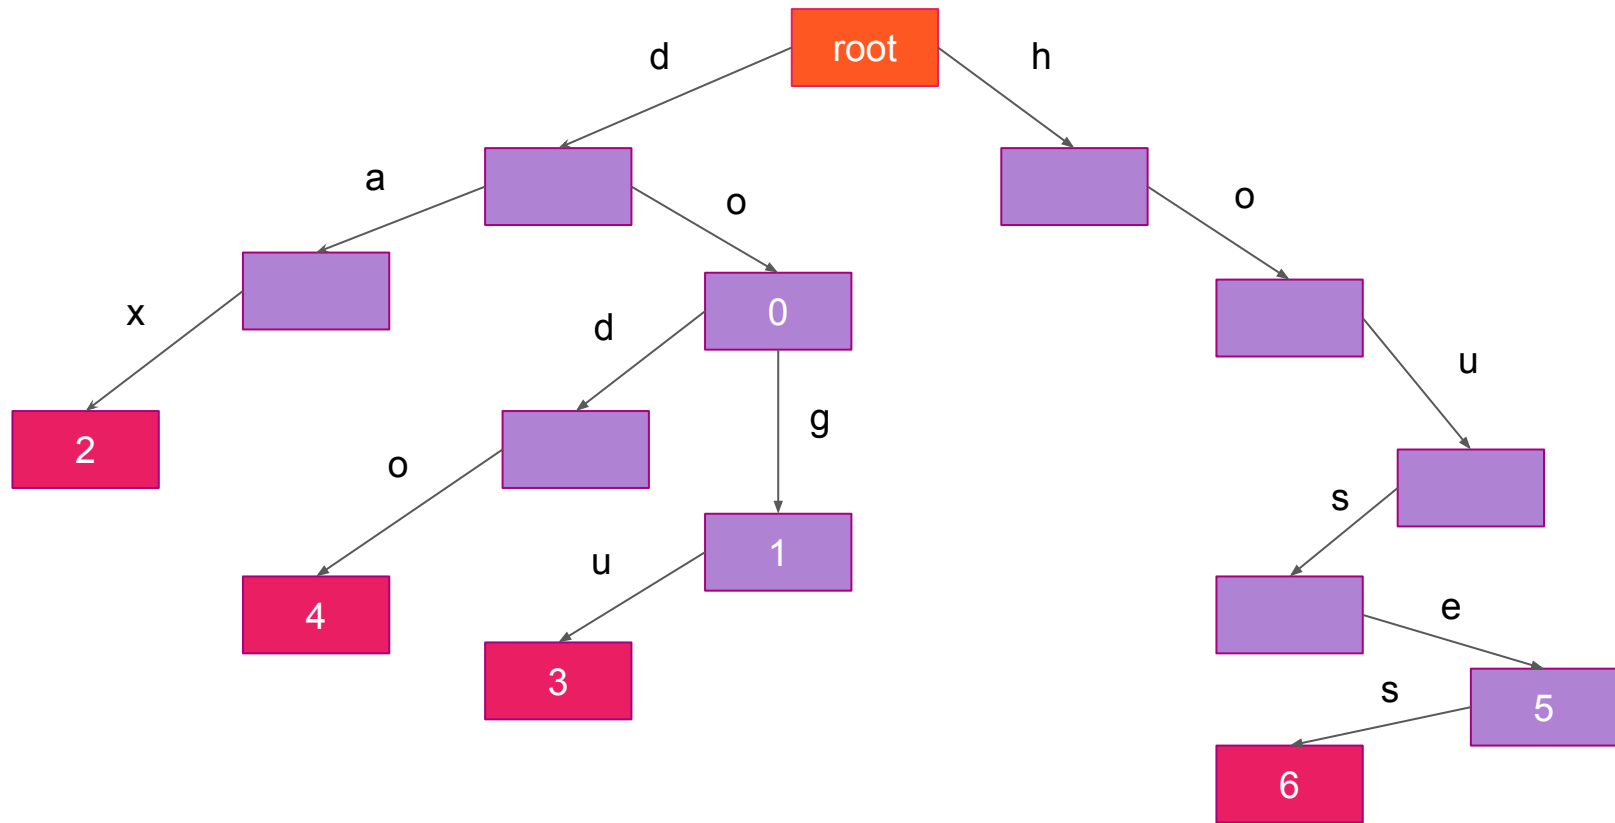
- Start at root
- Split string into characters
- For every character, follow an edge labelled by that character
- If edge does not exist, return false
- When you arrive at a node and your string is consumed, check if node is marked
- If it is marked, return **yes**
- Otherwise, return **no**

# Tries: example

**{ do: 0, dog: 1, dax: 2, dogu: 3, dodo: 4, house: 5, houses: 6 }**



# Tries




# Patricia (or radix) trie

- Space-optimized trie
- An isolated path (with nodes which are only children) with unmarked nodes is *merged* into one edge
- The label of the merged edge is the concatenation of the merged symbols

# Tries / Patricia tries as key/value store

- Marking does not need to be yes/no
- Can contain arbitrary value
- This allows us to map keys to values
- **add(key, value)**
- **query(key) → value**

# Patricia trie




# Merkle Patricia trie

- An authenticated Patricia Trie
- First implemented in Ethereum
- Allows proof-of-inclusion (of key, with particular value)
- Allows proof-of-non-inclusion (by showing key does not exist in trie)

# Merkle Patricia Trie

- Split nodes into three types:
  - **Leaf**: Stores edge string leading to it, and **value**
  - **Extension**: Stores **string** of a single edge, **pointer** to next node, and **value** if node marked
  - **Branch**: Stores one pointer to another node per alphabet symbol, and **value** if node marked
- We encode keys as hex, so alphabet size is 16
- We encode all child edges in every node with some encoding (e.g. JSON)
- Pointers are by hash application
- Arguments for correctness and security are same as for Merkle Trees

# Merkle Patricia trie



# Merkle patricia trie: node

key	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	value
-----	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	-------



# Merkle patricia trie: example


```
{ 'cab8': 'dog', 'cabe': 'cat', '39': 'chicken', '395': 'duck', '56f0': 'horse' }
```

root hash	0	1	2	$H_A$	4	$H_E$	6	7	8	9	A	B	$H_B$	D	E	F	
-----------	---	---	---	-------	---	-------	---	---	---	---	---	---	-------	---	---	---	--

$H_A$	9															$H_C$
-------	---	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-------

$H_B$	ab															$H_J$
-------	----	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-------

$H_C$	0	1	2	3	4	$H_D$	6	7	8	9	A	B	C	D	E	F	chicken
-------	---	---	---	---	---	-------	---	---	---	---	---	---	---	---	---	---	---------



Block Header,  $H$  or  $B_H$ stateRoot,  $H_r$ Keccak 256-bit hash of the root  
node of the state trie, after all  
transactions are executed and  
finalisations applied

Hash function:

KECCAK256()

World State Trie

Simplified World State,  $\sigma$ 

Keys

Values

a	7	1	1	3	5	5	45.0 ETH
a	7	7	d	3	3	7	1.00 WEI
a	7	f	9	3	6	5	1.1 ETH
a	7	7	d	3	9	7	0.12 ETH

ROOT: Extension Node

prefix	shared nibble(s)	next node
0	a7	

Branch Node

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value

Leaf Node

prefix	key-end	value
2	1355	45.0ETH

Extension Node

prefix	shared nibble(s)	next node
0	d3	

Leaf Node

prefix	key-end	value
2	9365	1.1ETH

Prefixes

0 - Extension Node,  
even number of nibbles  
1□ - Extension Node,  
odd number of nibbles,  
2 - Leaf Node, even  
number of nibbles  
3□ - Leaf Node, odd  
number of nibbles  
□ = 1<sup>st</sup> nibble  
1 nibble = 4 bits

Branch Node

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value


Leaf Node

prefix	key-end	value
3□	7	1.00WEI

Leaf Node

prefix	key-end	value
3□	7	0.12ETH

# Blocks



- Data structure with three parts:
  - nonce (ctr), data (x), reference (s)
  - Typically called the **block header**
- data (x) is application-dependent
  - In Bitcoin it stores financial data (UTXO-based)
  - In Ethereum it stores contract data (account-based)
  - In Namecoin it stores name data
  - We leave this undefined for now -- we will come back to this in future lectures
- Block validity:
  - Data must be valid (application-defined validity)

# Proof-of-work in blocks


- Blocks must satisfy proof-of-work equation

$$H(\text{ctr} \parallel \mathbf{x} \parallel \text{s}) \leq T$$

- for some constant  $T$
- ctr is the nonce used to solve proof-of-work
- The value  $H(\text{ctr} \parallel \mathbf{x} \parallel \text{s})$  is known as the **blockid**

# Blockchain


- Each block references a **previous** block
- This reference is by **hash** to its **previous** block, similar to Merkle Trees
- This linked list is called the **blockchain**
- Convention: Arrows show authenticated inclusion



- Blocks use the `s` value to point to the previous block by hash

# Blockchain


- The **first** block of a blockchain is called the Genesis Block






# Notation conventions

- We use the symbol **C** to denote a blockchain
- C is a **sequence** of blocks
- We use  $C[i]$  to denote the  $i^{\text{th}}$  block (0-based)
- $C[0]$  denotes genesis
- We use  $C[-i]$  to denote the  $i^{\text{th}}$  block from the end
- Chain property: For each  $i > 0$ :  $C[i].s = H(C[i - 1])$



# Notation conventions

- Range notation:  $C[i:j]$  denotes a **subsequence** from  $i$  (inclusive) to  $j$  (exclusive)
- Similarly  $C[-i:j]$ ,  $C[i:-j]$ ,  $C[-i:-j]$




# Notation conventions

- Range notation:


$C[i:]$ ,  $C[-i:]$  denotes chain from  $i$  (or  $-i$ ) inclusive to end inclusive


$C[:i]$ ,  $C[:-i]$  denotes chain from beginning inclusive to  $i$  (or  $-i$ ) exclusive

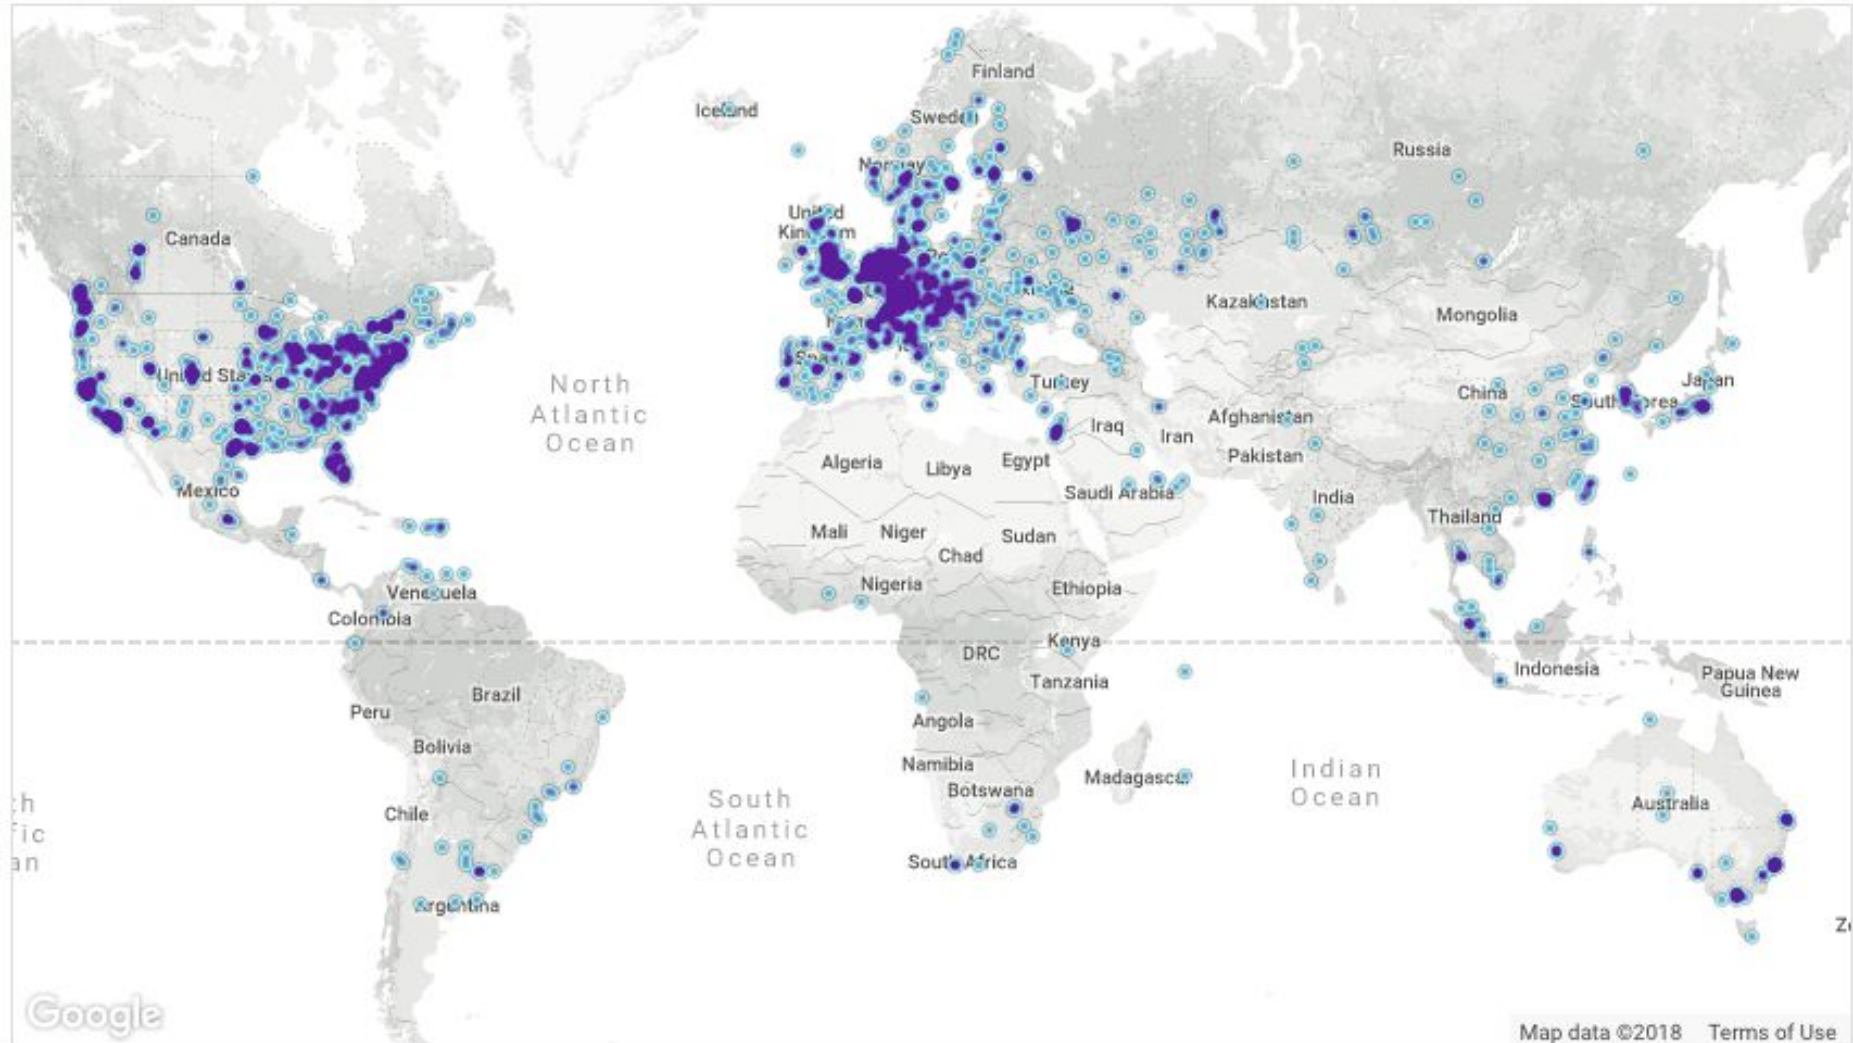


# The bitcoin network

- All bitcoin nodes connect to a common p2p network
- Each node runs the code of bitcoin
- A node can run on a phone, computer, etc.
- Open source code
- Each node connects to its neighbours
- They continuously exchange financial data
- Each node can **freely** enter the network -- no permission needed! A “permissionless network”.
- **The adversarial assumption:**  
There is no trust on the network! Each neighbour can lie.







# Peer discovery

- Each node stores a list of peers (by IP address)
- When Alice connects to Bob, Bob sends Alice his own known peers
- That way, Alice can learn about new peers



# Bootstrapping the p2p network

- Peer-to-peer nodes come “preinstalled” with some peers by IP / host
- When running a node, you can specify extra “known peers”

# The *gossip* protocol

- When a node **Alice** generates some new data...
- Alice **broadcasts** data to its peers
- Each peer broadcasts this data to *its* peers
- If a peer has seen this data before, it ignores it
- If this data is new, it broadcasts it to its peers
- That way, the data spreads like an epidemic, until the whole network learns it
- This process is called **diffuse**

# Financial data

- Financial data is encoded in the form of *transactions*
- Every transaction is broadcast on the network to everyone using the gossip protocol
- Financial data on cryptocurrencies are **common knowledge** among all participants

Transactions on [blockchain.info](https://blockchain.info)

# Eclipse attacks

- Isolate some honest nodes in the network effectively causing a “net split” in two partitions A and B
- If peers in A and peers in B are disjoint and don’t know about each other, the networks will remain isolated
  - More recent attack: Erebus
- The connectivity assumption:  
There is a path between two nodes on the network  
**If a node broadcasts a message, every other node *will* learn it**

