

Blockchains & Distributed Ledgers

Lecture 08

Aggelos Kiayias

Slide credits: AK

Lecture Overview

- Anonymity & Privacy in blockchain protocols.
 - Bitcoin and CoinJoin transactions.
 - Mix-nets
 - group and ring signatures.
 - Cryptonote/Monero
 - Zero-knowledge proofs & SNARKs
 - Zcash.

Pseudonymity vs. Anonymity


- Pseudonymity: identities are substituted by tags that are independently assigned to each identity.
- Anonymity: any action performed is manifested within a set of indistinguishably acting participants.
(The anonymity set)

Privacy and Bitcoin

- Users can create accounts -practically- without cost and without association to previous accounts.
- Essentially they can create an unlimited number of pseudonyms.


Transaction Graph Analysis

blockchain





coinbase
transaction

account *a*
moves 50 BTC to
accounts *b* and *c*
(minus fees)




Common Behaviours



Fungibility and Privacy

- Coins are interchangeable.
- Since each “satoshi” has its whole history in the bitcoin blockchain, its fungibility is debatable.

Anonymising Transactions




Anonymity
set of size n


TP may disappear
with the money

CoinJoin

Without CoinJoin



With CoinJoin



Using Multiple Input Transactions

- parties broadcast a_1', a_2', \dots, a_n' accounts to each other.
- The i -th party broadcasts a signature from account a_i to pay the account a_i' from the set of accounts a_1', a_2', \dots, a_n' . When all n signatures are broadcasted then the multiple input transaction can be posted on the blockchain.
- If any of the n parties abort the protocol the transaction cannot be validated.
- **Questions:** how to ensure that the adversary cannot do a correlation between a_i and a_i' ? in case of an abort how to restart the protocol without the offending party?


Passive vs Active Attacks

- A “passive” adversary would just observe the transaction in the blockchain. In this case, an anonymity set of size n protects each participant.
- However, an “active” adversary participates in a protocol execution; the correlation between participants is apparent due to the broadcast.

Mix-net

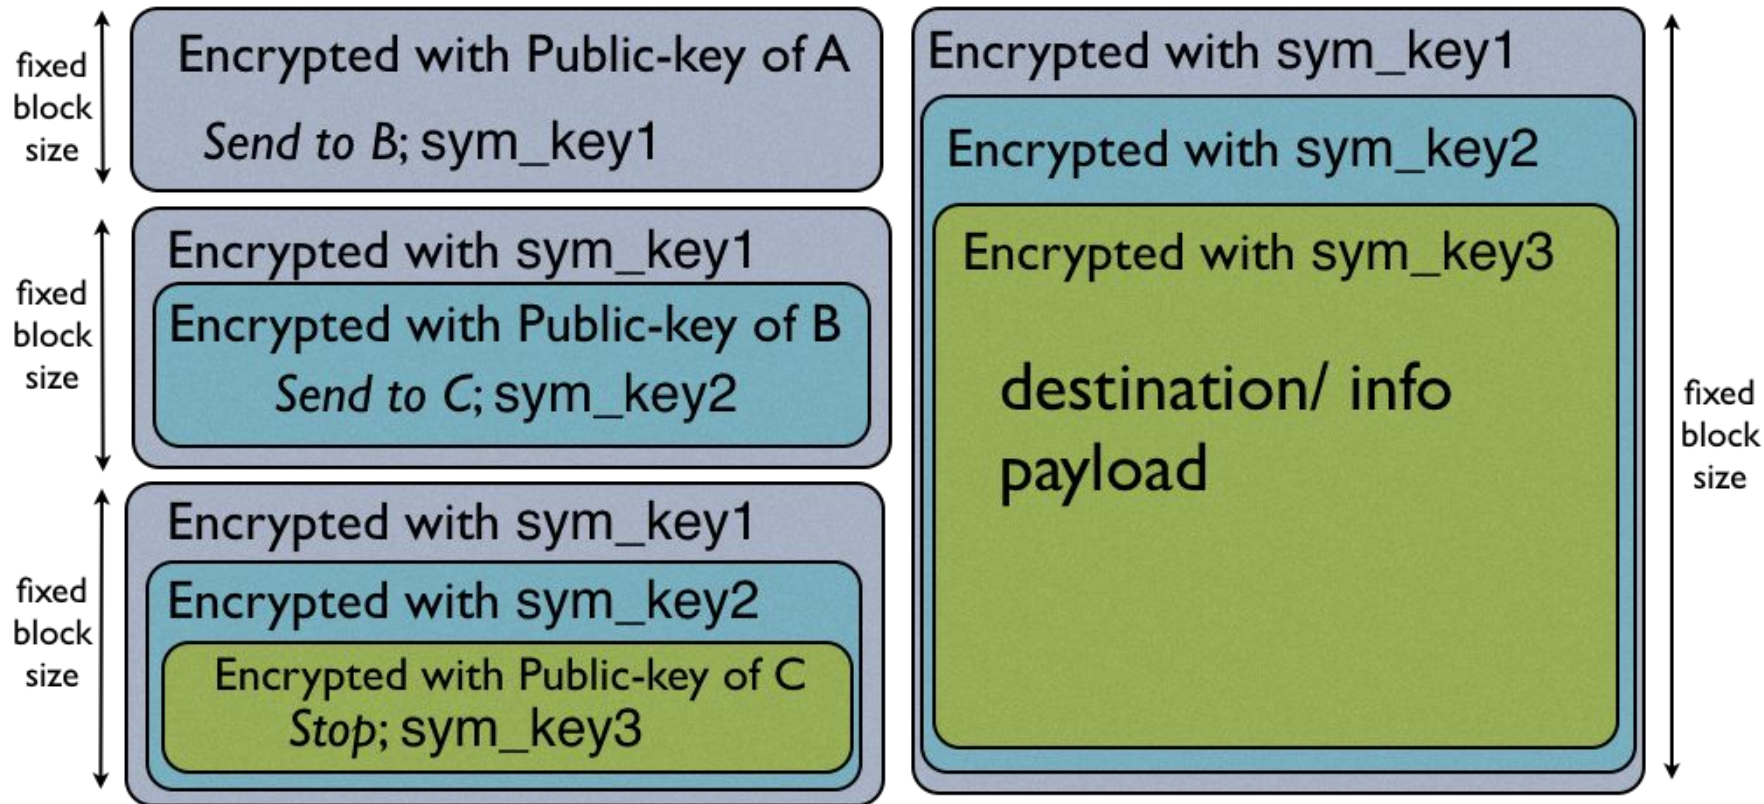
- Facilitates a sender-anonymous broadcast.
 - Decryption mix-nets.
 - Re-encryption mix-nets.

Mix-net, 2




Not possible to relate whether P send msg1 or msg2
and similarly for Q (as long as there is one honest mix)

Decryption Mix-net



Routing via a Mix-net




Mix-net for Coinjoin transactions

- Parties broadcast their public-keys (the association between public-keys and accounts a_1, a_2, \dots, a_n is public).
- Parties engage in a decryption mix-net in sequence so that the last party P_n obtains the sequence of accounts a_1', a_2', \dots, a_n' . P_n broadcasts the accounts to all.
- Note that each step is performed by a designated party P_i , hence any abort can be attributed to that party. A repeat session may exclude the party P_i .


Hiding Coin Balances

balances
are visible:

Without CoinJoin



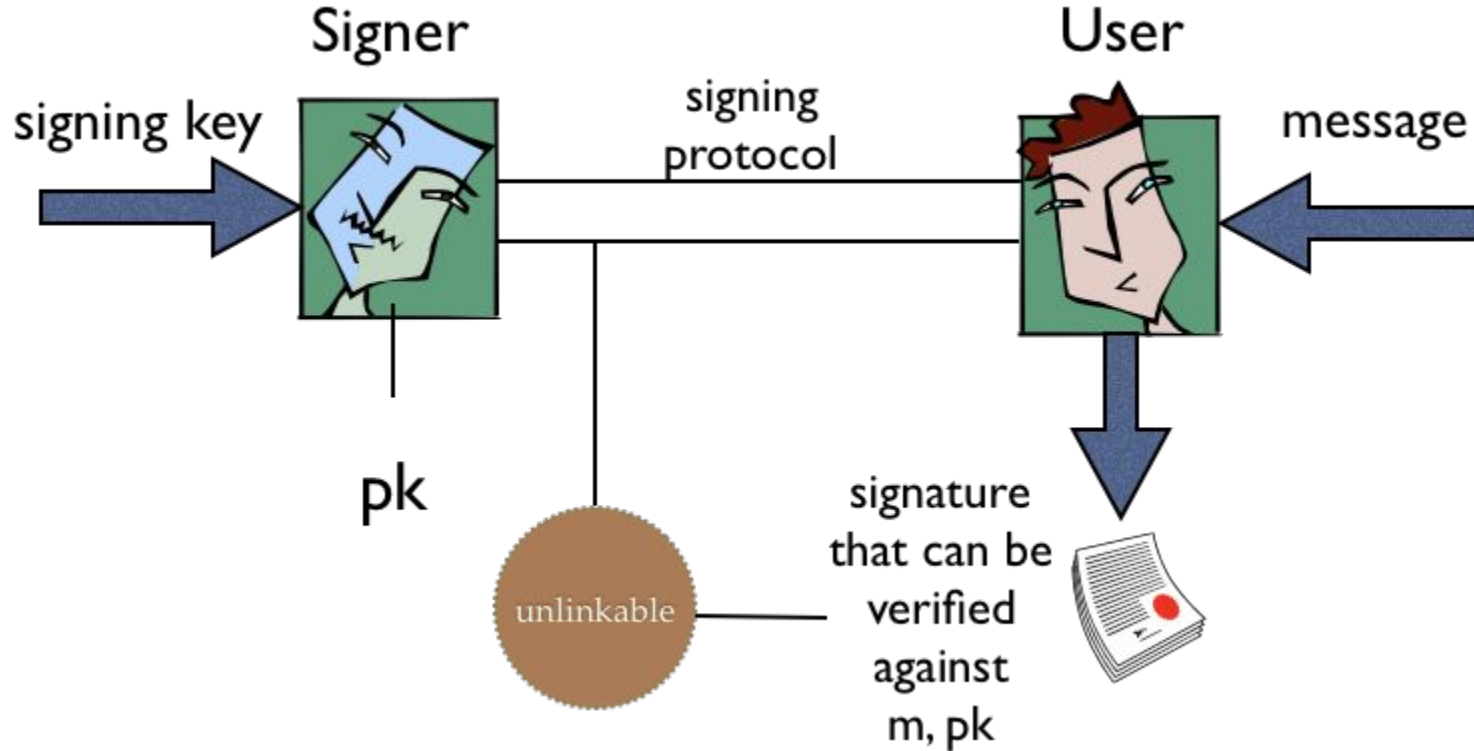
With CoinJoin



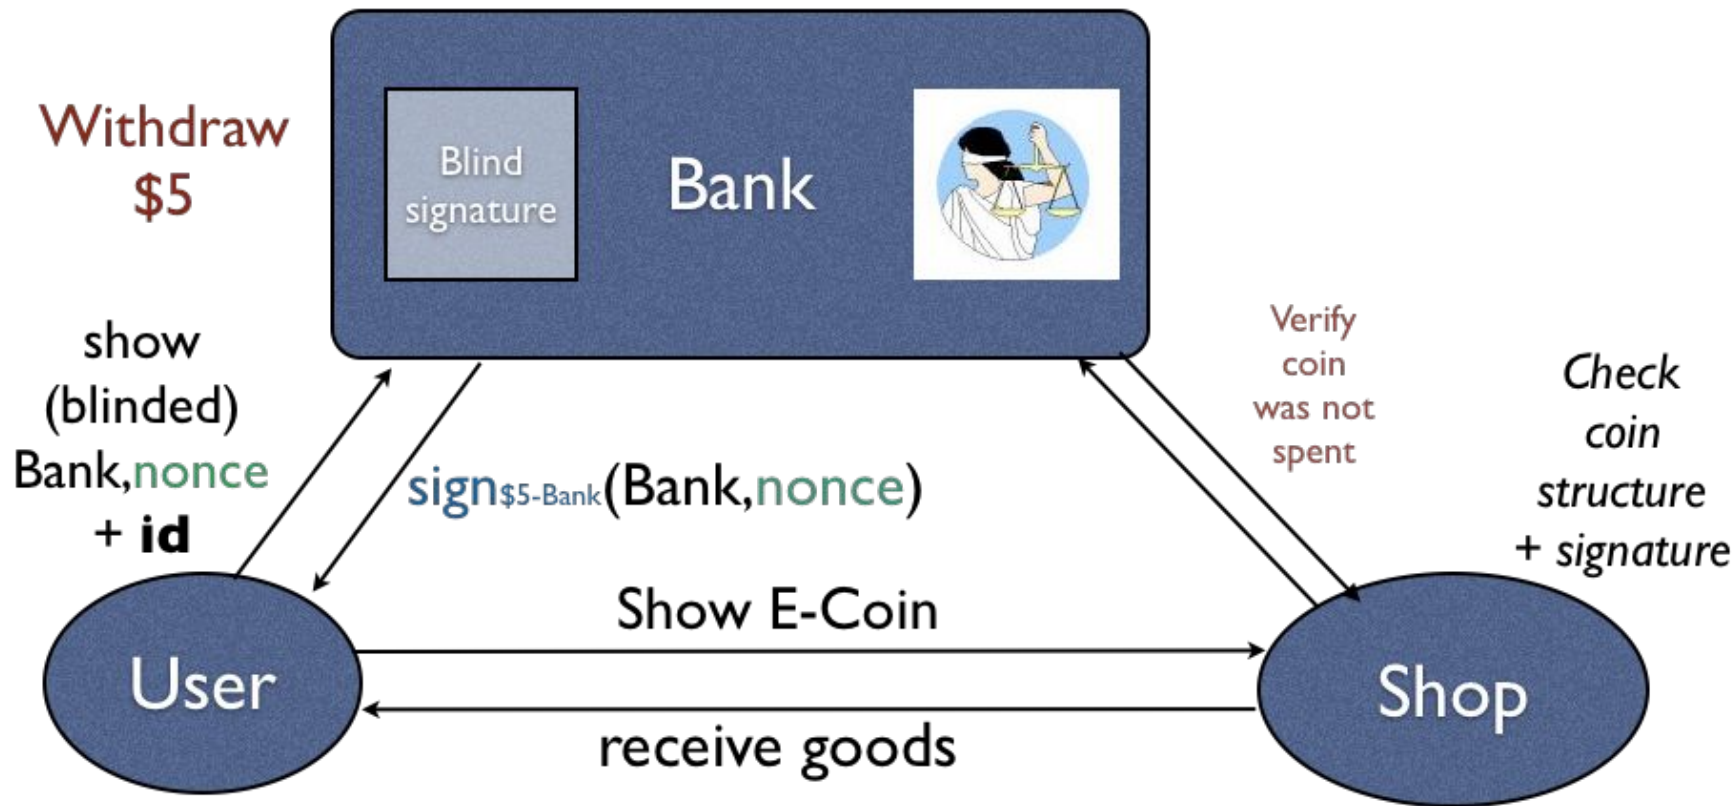
Coordination

- Coinjoin & similar techniques require coordination and message passing between the parties engaged in the transaction.
- Is it possible to improve that?
 - a. Using an intermediate party.
 - b. Using more advanced cryptographic techniques.

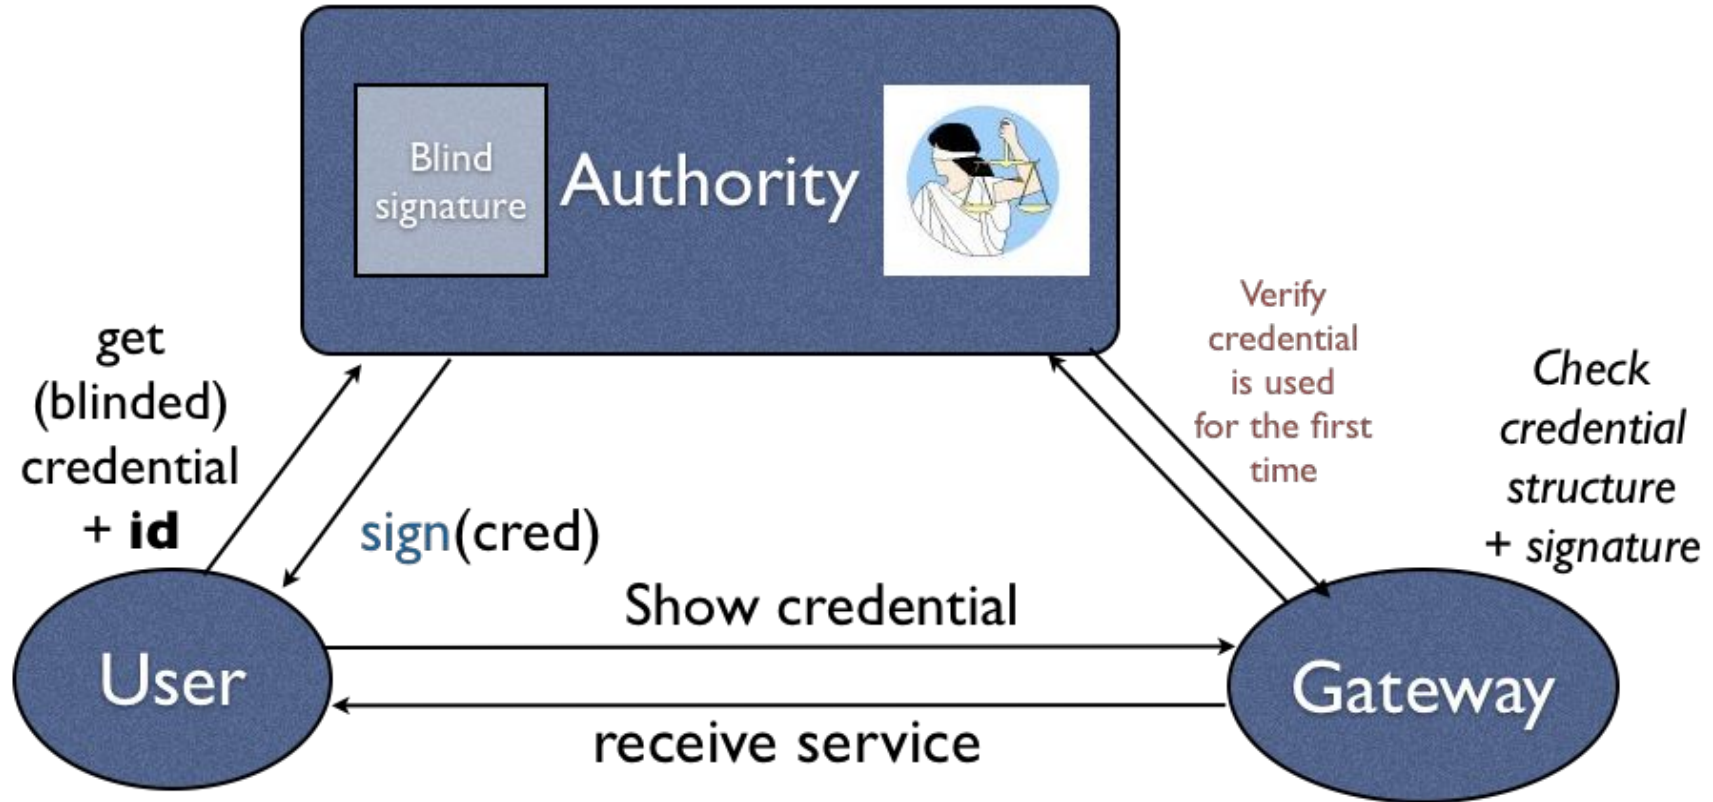
Blind-Signatures



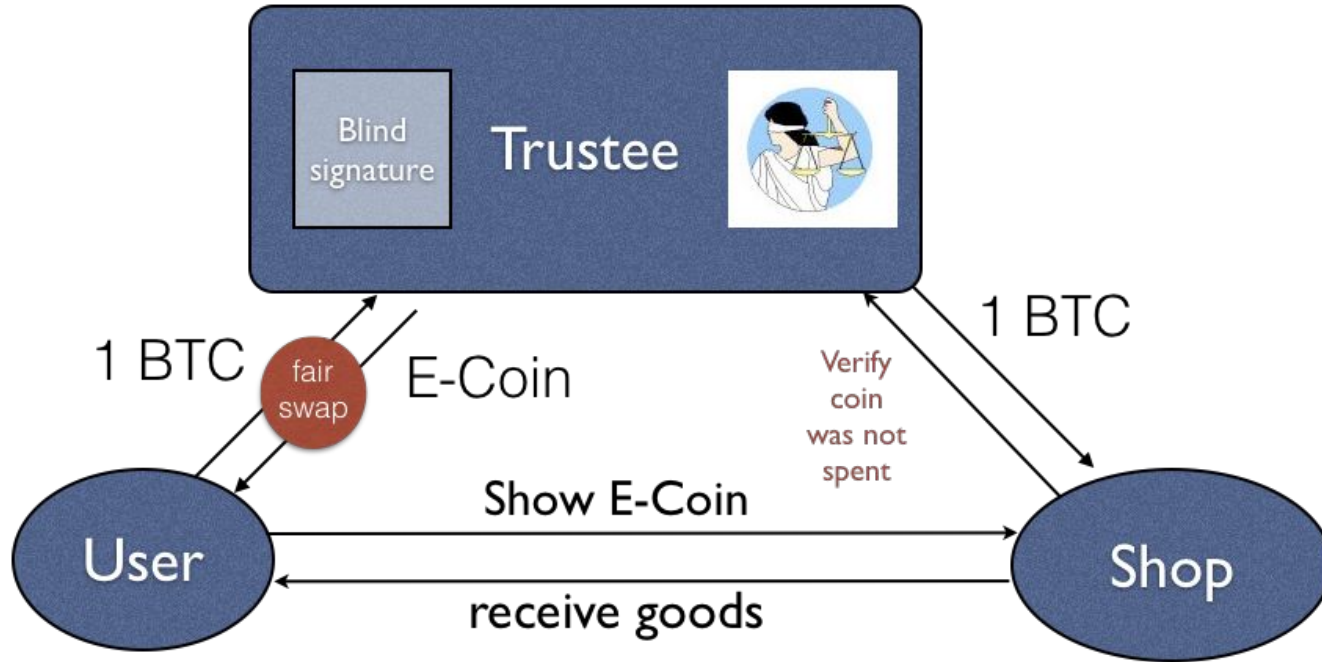
Chaum's E-cash



Anonymous Credentials



Anonymizing Bitcoin Payments via E-cash



Note: Trustee is trusted to honor its e-coins.

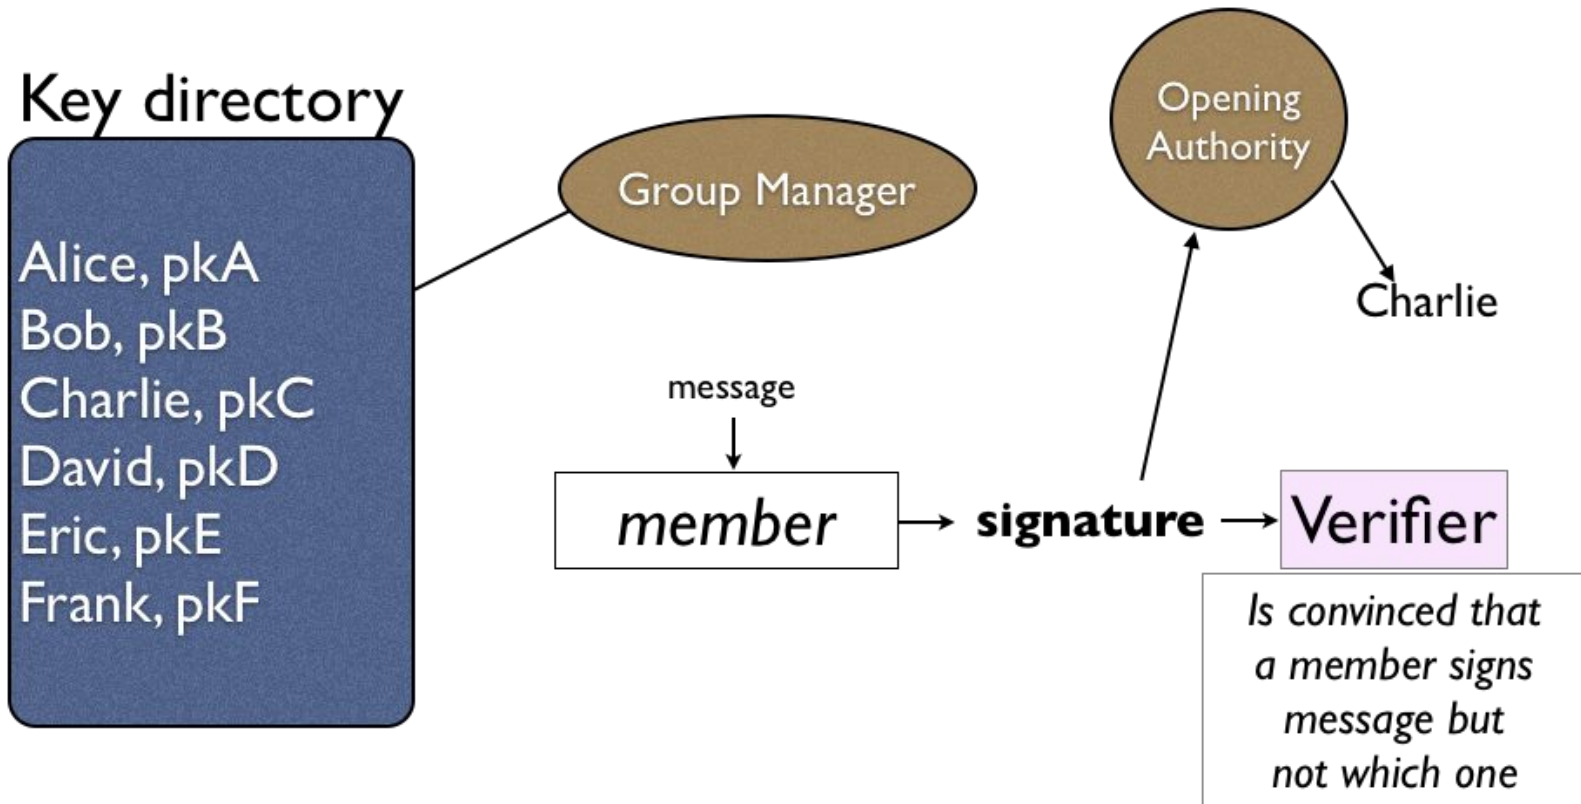
Fair Swaps

- Alice and Bob would like to exchange secrets so that either none of them gets output or both.
 - Classical problem!
 - Impossible to solve under standard network assumptions!
 - Going around the impossibility: (i) optimistic fair exchange (ii) resource-based fair exchange (iii) fair swaps with penalties.
- [Construction] using a smart contract that both parties fund to accept their secrets. Key requirements: (i) parties lock up funds, (ii) secret submission can be verified by the contract code

Anonymity and Digital Signatures


- So far all digital signatures identify the signer.
 - Is it possible to hide the sender within a group?

Group Signatures



Traceable Signatures

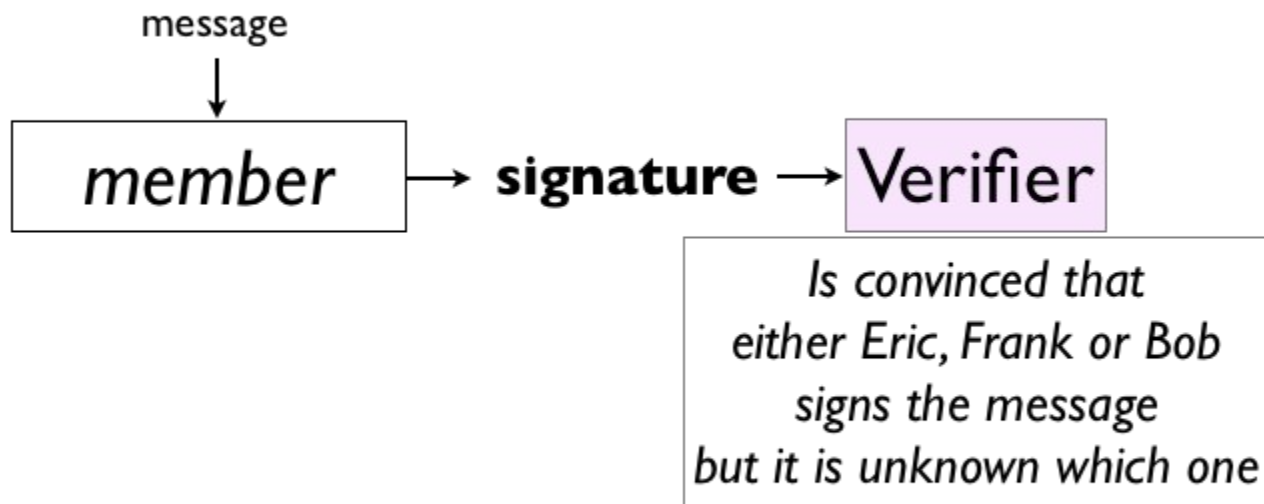
Traceable Signatures



Ring Signatures

Key directory

Alice, pkA
Bob, pkB
Charlie, pkC
David, pkD
Eric, pkE
Frank, pkF



Monero/Cryptonote

- Uses “*stealth*” addresses and *linkable ring signatures* to provide better anonymity.
 - a. For each payment an anonymity set is selected with accounts of the same monetary value.
 - b. A ring signature is issued on behalf of that set, suitably restricted so that an account can only be used once (if twice it is *linkable*).
 - c. Stealth addresses enable the sender to create unlinkable addresses for the receiver which subsequently the receiver can detect.

Is Monero Anonymous?

- There is potentially more uncertainty in terms of transactions compared to a bitcoin-like blockchain.
- Nevertheless, it is not obvious how to quantify the level of anonymization.
- De-anonymization is feasible in a number of cases. (e.g., imagine the attacker “spraying” the ledger with transactions so that it commands a good number of selected accounts)

Increasing and Safeguarding the anonymity set, I

- A larger anonymity set is most preferable.
- However in the techniques we have seen so far, transaction preparation work increases linearly with the anonymity set.
- **Ideal:** use the set of all possible unspent transaction outputs.

Increasing the anonymity set, II

$$\langle \rho, sn, \psi = \underbrace{\text{Commit}(\rho, sn)}_{\text{public}} \rangle$$

The commitment value is associated with a deposit to the ledger (“minting” a coin for \$1).

Spending a coin, requires announcing the sn and proving that it was committed before in the ledger; (withdrawing \$1)

$$\underbrace{\exists i : \psi_i = \text{Commit}(\rho, sn)}$$

existential quantifier over all commitments in the blockchain

Increasing the anonymity set, III

Organize all commitments and serial numbers in a Merkle tree.

Prove that there is a leaf in the Merkle tree that contains the commitment

$$\psi_i = \text{Commit}(\rho, sn)$$

Statement representation and witness size logarithmic in the number of coins.

Challenges

- How is it possible to prove efficiently statement referring to the leaf of a Merkle tree?
 - a possible solution: use “ZK-snarks”
- Transferring a coin from one user to another is not properly specified (one cannot simply transfer ρ).

ZK-Snarks

- Zero-knowledge succinct arguments of knowledge.
 - like zero-knowledge proofs, but with:
 - *computational* soundness: it depends on the security of a “common reference string”; this is structured cryptographic information that is assumed to be honestly sampled.
 - succinctness: the proof size and the verifier’s running time is efficient proportionally to the *statement* only.

Constructing ZK-SNARKs

- There exist a SNARK for any NP-relation R .
- The actual proof sizes are small (hundreds of bytes)
- Verification does not depend on the running time of R .

Zerocash

$$\langle a_{\mathbf{pk}}, \underbrace{v}_{\text{value}}, \underbrace{s}_{\text{random}} \rangle \quad (a_{\mathbf{pk}}, a_{\mathbf{sk}})$$

account
public/secret
key

$$k = \text{Commit}(\rho, a_{\mathbf{pk}} || s)$$

$$sn = \text{PRF}_{a_{\mathbf{sk}}}^{sn}(s)$$

$$\psi = \text{Commit}(\rho', v || k)$$

$$\text{coin} : \langle a_{\mathbf{pk}}, v, s, \rho, \rho', \psi \rangle$$

The double commitment enables verifying that the value v is properly encoded in the coin without revealing information about the owner

ZeroCash “Pour” Operation

given coin $\langle a_{\text{pk}}, v, s, \rho, \rho', \psi \rangle$

produce two new coins with values $v_1 + v_2 = v$

$$a_{\text{pk}}^1, a_{\text{pk}}^2$$

Serial number of spent coin is revealed and marked as spent

set $k_i = \text{Commit}(\rho_i, a_{\text{pk}}^i || s_i)$

$$\psi_i = \text{Commit}(\rho'_i, v_i || k_i)$$

Reveal ψ_1, ψ_2 and prove that the Merkle tree has a commitment corresponding to a coin

$\langle a_{\text{pk}}, v, s, \rho, \rho', \psi \rangle$ that is split properly and a_{sk} is known

Include a public-key encryption of opening the commitment that the recipient can use to decrypt the coin secret values.

How to obtain a CRS?

- Trusted computation is needed.
 - a. Use *secure multiparty computation* (topic of our next lecture)
 - b. Use *updateable reference strings* (URS) instead and outsource the update operation.
 - c. Use *alternatives* to SNARKs that do not require it (disadvantage: performance would be worse)