

Blockchains and Distributed Ledgers Lecture 01

Aggelos Kiayias



THE UNIVERSITY
of EDINBURGH

Lecture 01

- Introduction to blockchain.
- What is a distributed ledger.
- Hash functions and digital signatures.
- Transactions in Bitcoin.

Understanding DL's

- Distributed ledgers use blockchain protocols as one main means of implementation.
- The blockchain is a distributed database that satisfies a unique set of safety and liveness properties.
- To understand it, we can focus to its first (and so far most successful) application.

Case study : Money

- What is money?

Money is useful

(1874) A man offering chicken for a yearly newspaper subscription



Properties of Money

can be used as medium
for the exchange
of goods - no bartering

- a medium of exchange
- a unit of account
- a store of value

can be used for
pricing of all goods
and services, for
accounting purposes
and debt recording.

storing and retrieving it at a
point in the future maintains
its value.

Creating Money

Money 1.0 : using a trusted object



Analysis of Money 1.0

mediocre

[ok for face to face transactions]

- a medium of exchange
- a unit of account
- a store of value

mediocre fungible,
but not divisible well.
typically forgeable.

bad. some objects may
deteriorate, others may have
unknown hidden quantities.

Creating Money

Money 2.0 : using a trusted entity



Trusted entity issues “IOU”s

Analysis of Money 2.0

good

[for transactions
within the domain of
the trusted entity]

- a medium of exchange
- a unit of account
- a store of value

great!

fungible & divisible.

mediocre

[tied to the availability & reputation
of the issuing entity]
[possible

Creating Money


Money 3.0 : Bitcoin

Enter Blockchains & Distributed Ledgers

The never-ending book parable



A “book” of transactions




- Each new page requires some effort to produce.
- Anyone can be a scribe and produce a page.
- New pages are produced indefinitely as long as scribes are interested in doing so.

Importance of Consensus

- If multiple conflicting books exist, which is the “right one”?

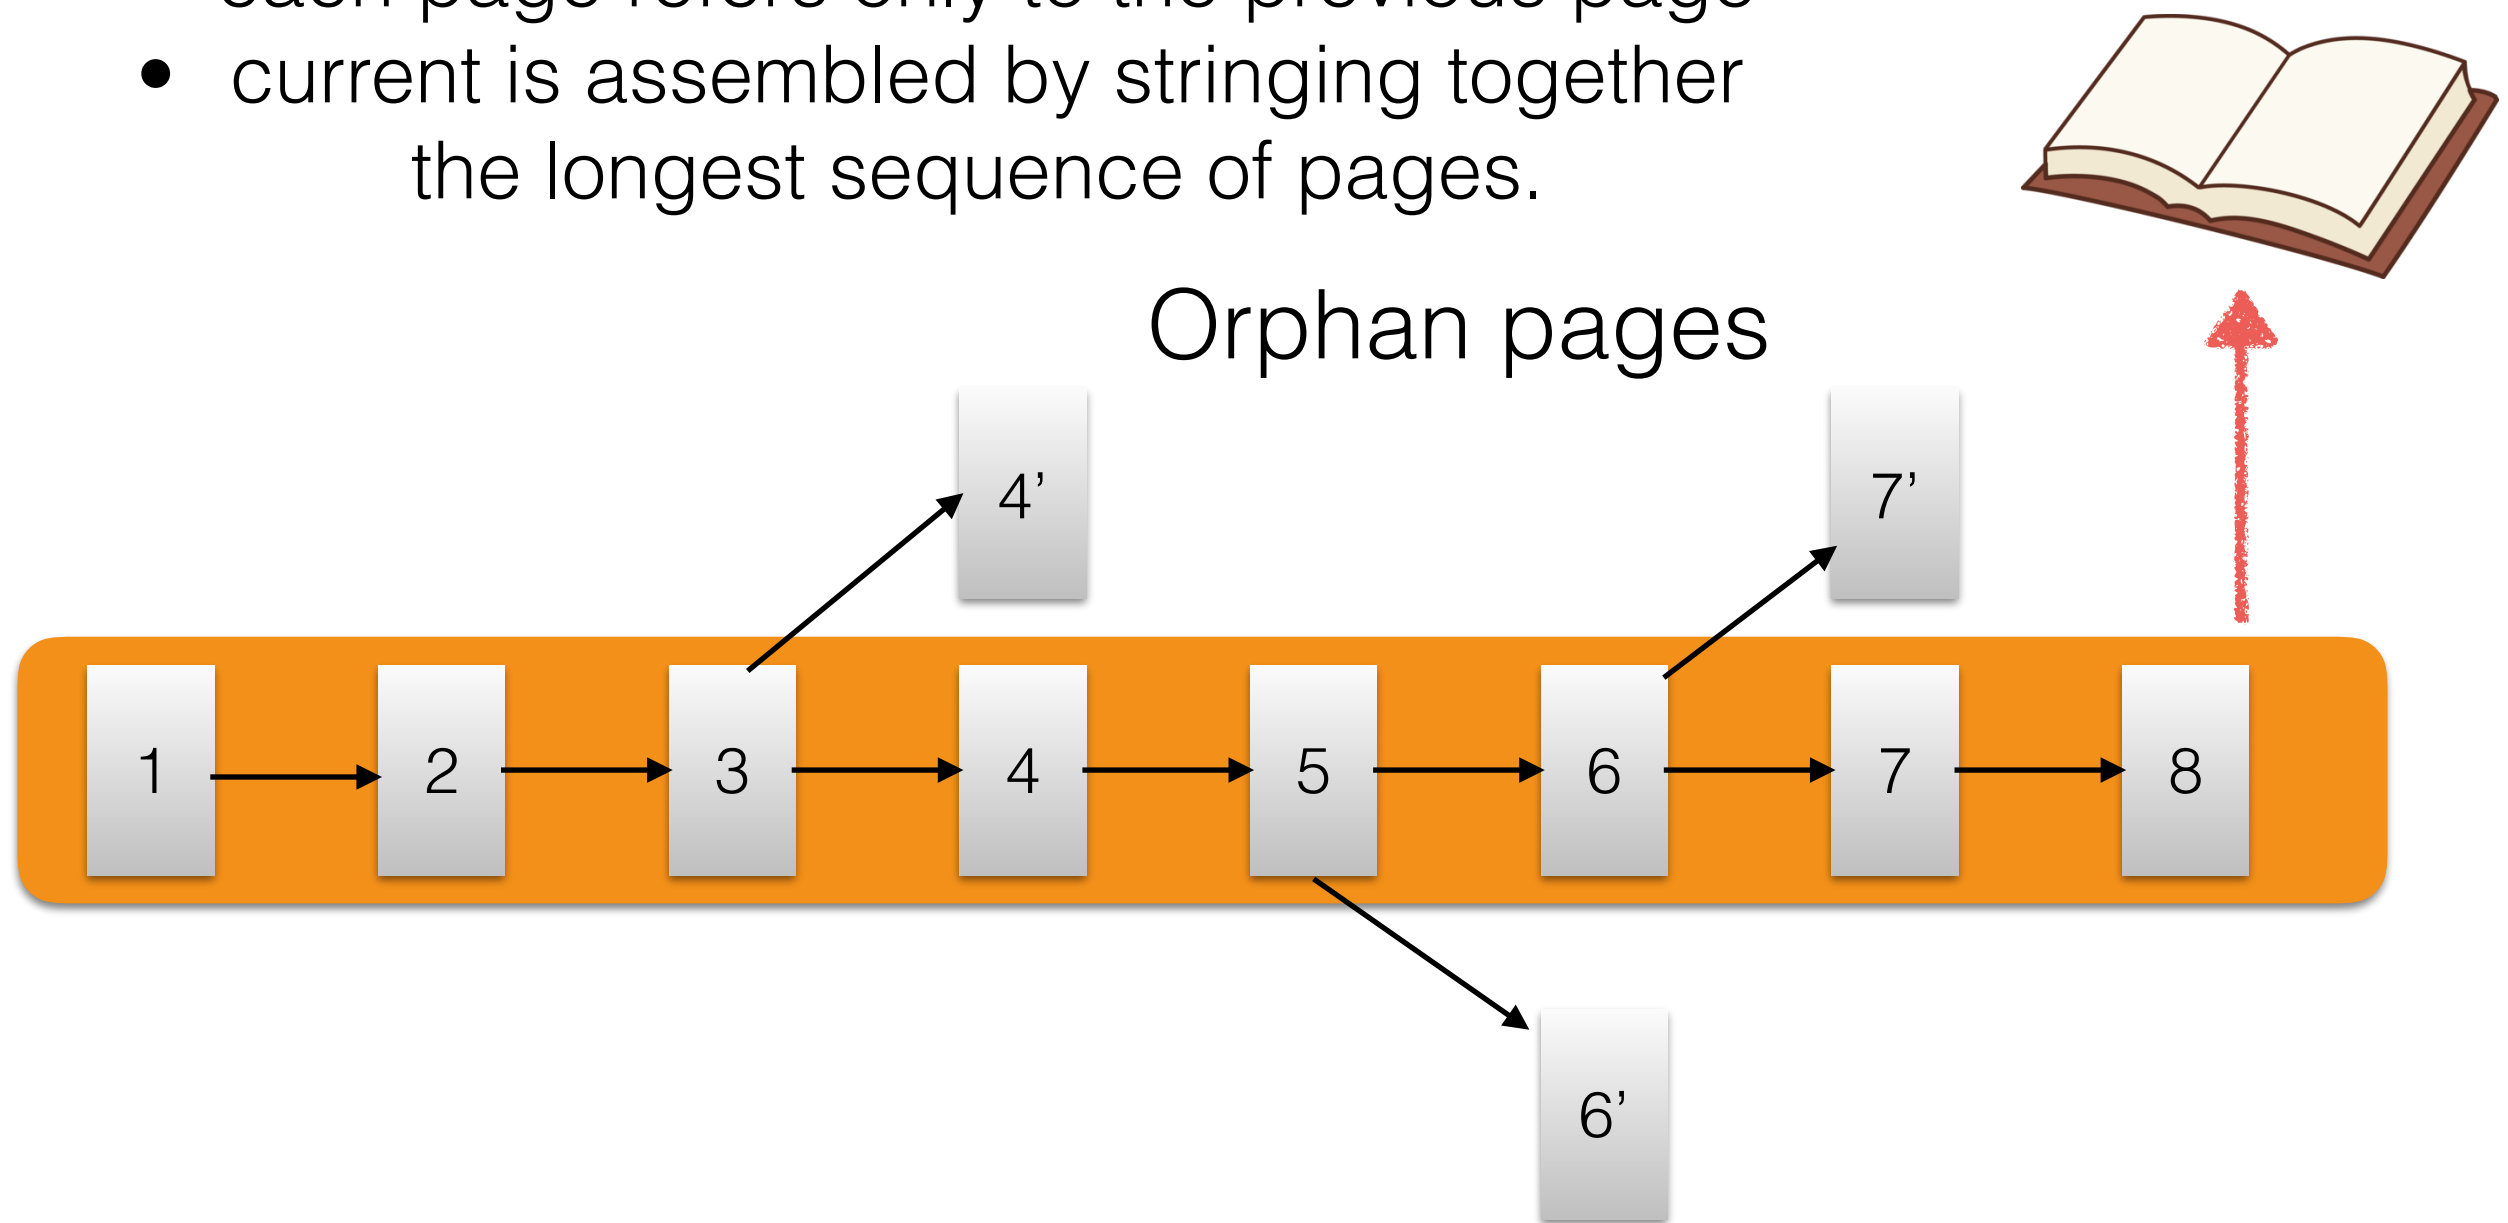
Choosing the correct book



The **current book** to work on & refer to is the book with the most pages. if multiple exist, just pick one at random.

Assembling the current book

- each page refers only to the previous page
- current is assembled by stringing together the longest sequence of pages.



Rules of extending the book





The first scribe that discovers
a page announces it to everyone else



Effort is needed to produce a page

equivalent to : each page needs a special
combination from a set of dice to be rolled.



The probabilistic
nature of the process
is paramount to its
security


The benefits of randomness

Imagine two scribes

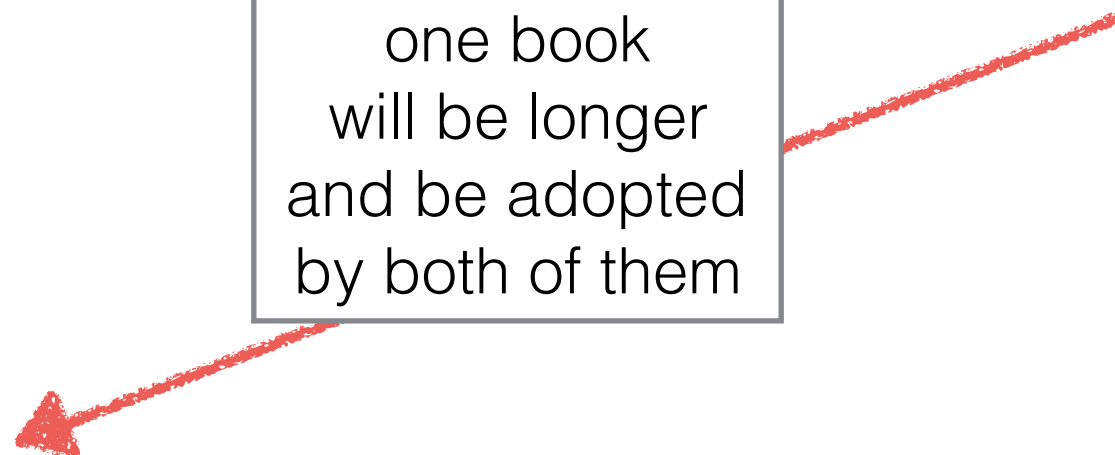

working together



Unlikely
to continuously
be lucky
together



eventually
one book
will be longer
and be adopted
by both of them

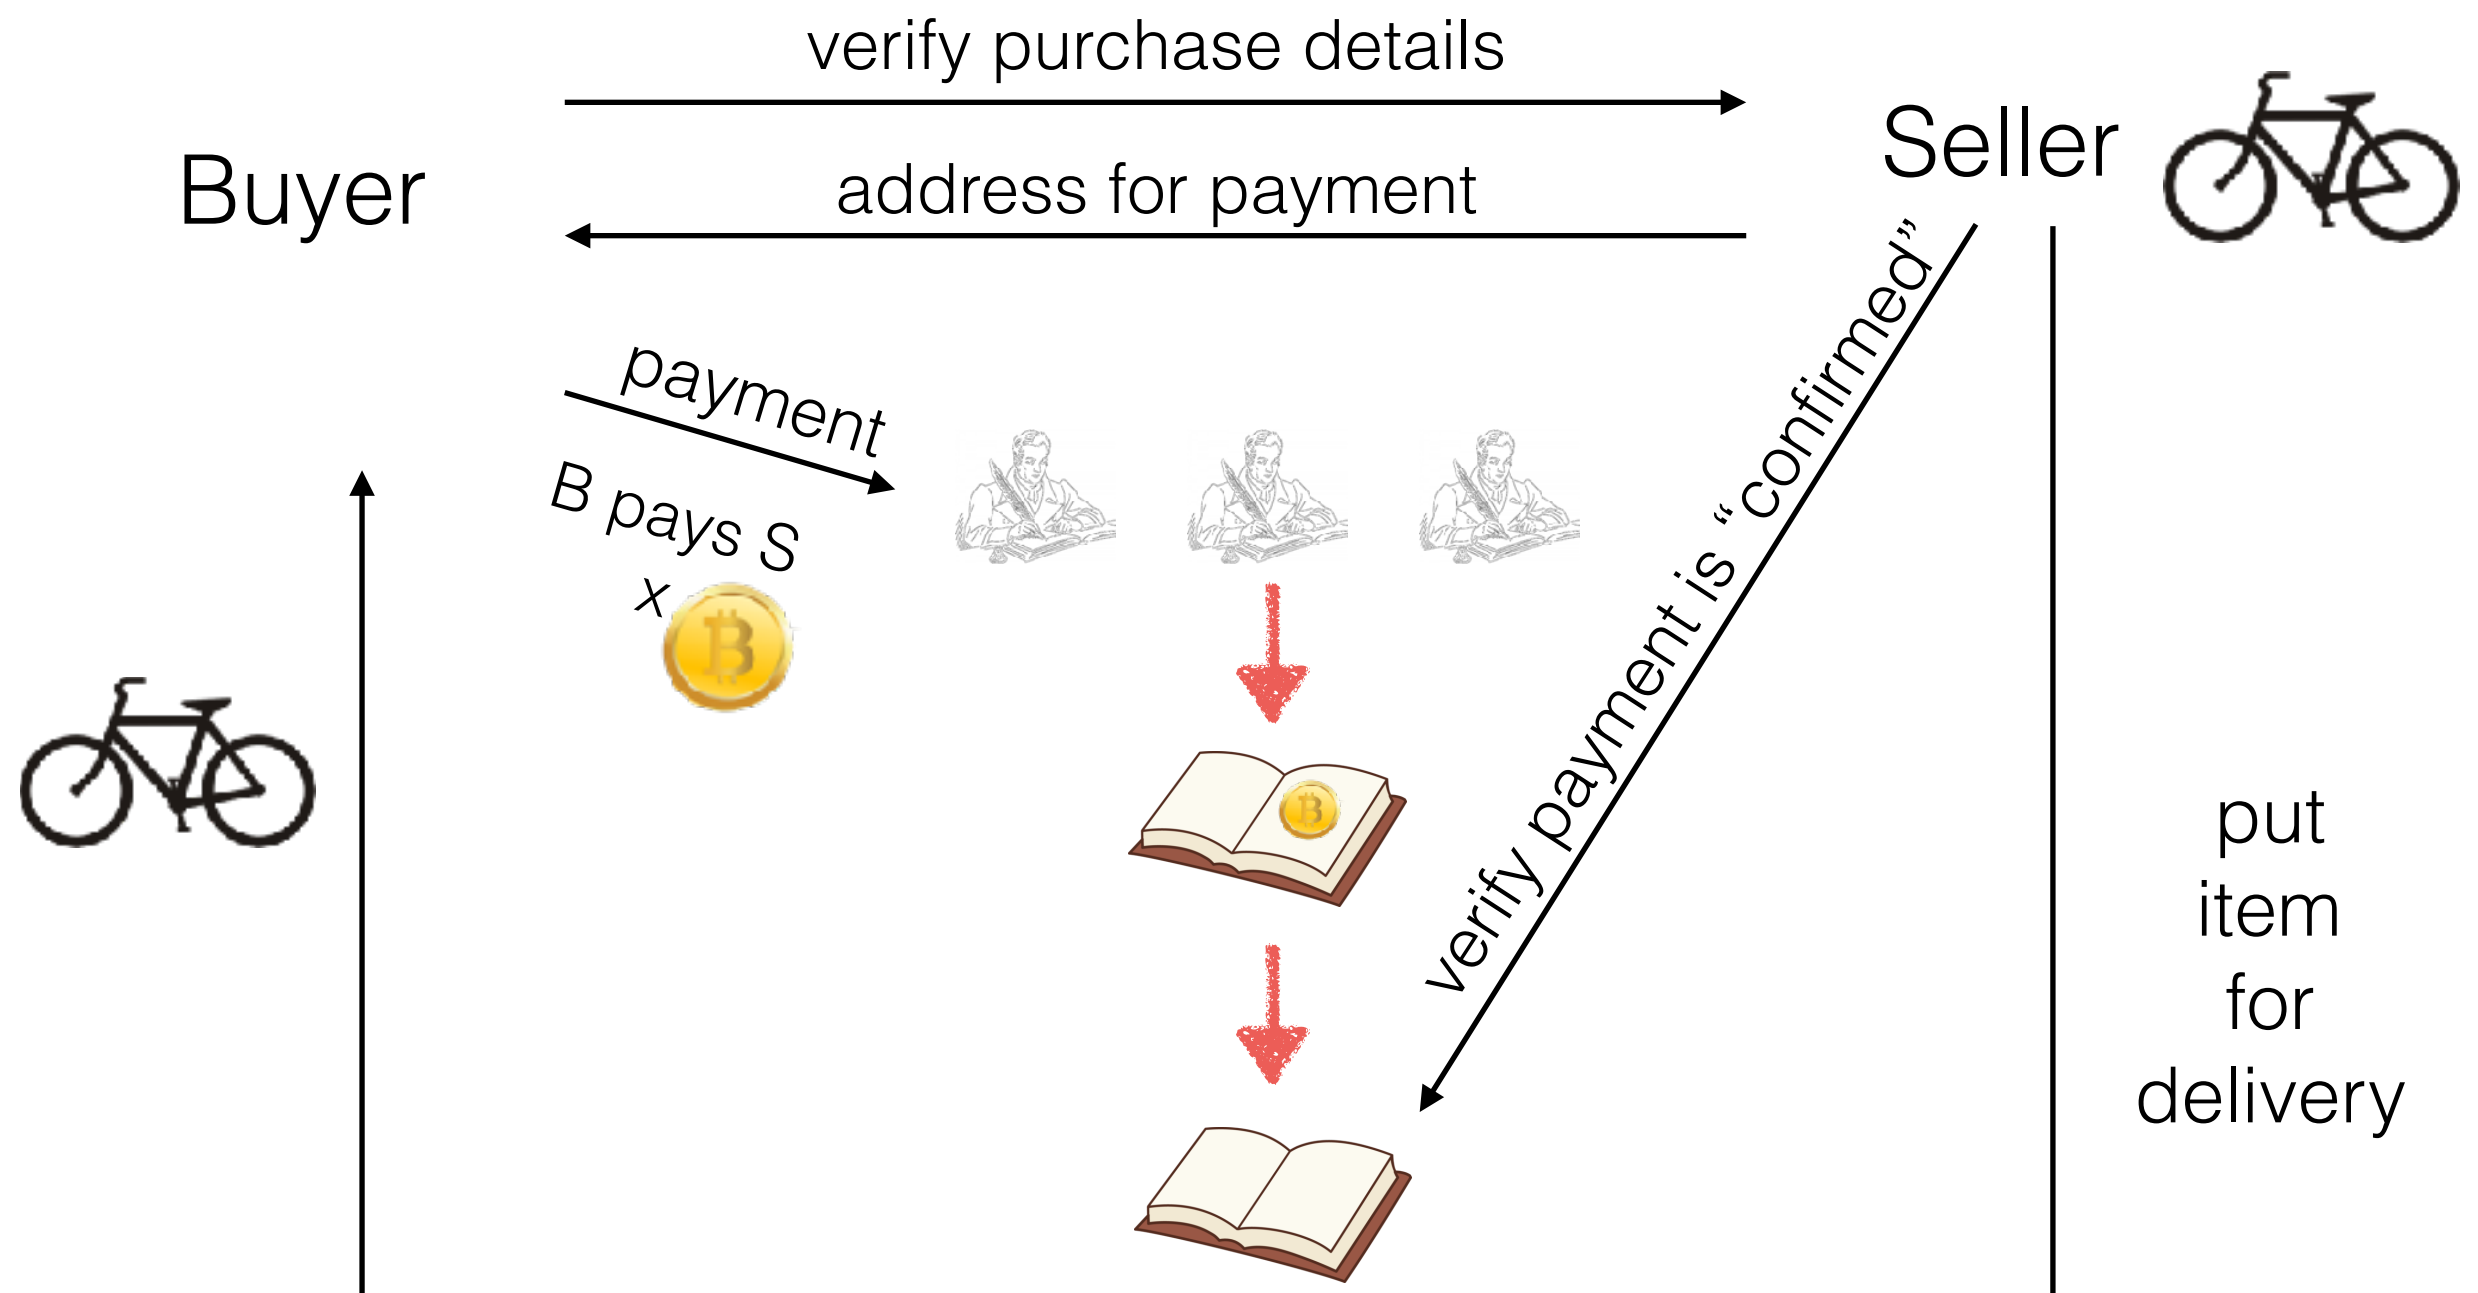


Symmetry Breaking

Being a scribe

- Anyone can be a scribe for the book.
- As long as one has a set of dice.
- The more dice one has, the higher the likelihood to produce the winning combination to make a page.

Using the book - Money 3.0



Parable & Reality

book




the “blockchain”

scribes



“Miners” / Computer systems
that organize
transactions in blocks

producing
a page



Solving a **cryptographic
puzzle** that is **moderately
hard** to solve

rolling
a set of
dice



Using a computer to test for a solution
from a large space of candidate
solutions

Analysis of Money 3.0

improving

[assuming internet
connectivity / adoption]

- a medium of exchange
- a unit of account
- a store of value

great!

fungible & divisible.

good

[no trusted parties -
no natural deterioration]

Word of Caution

- Just because something **can** be good as a store of value, it does **not** mean that it **will be** a good store of value in a real world deployment.



Word of Caution

- Just because something **can** be good as a store of value, it does **not** mean that it **will be** a good store of value in a real world deployment.

Dogecoin Charts



From Money to Smart Contracts

- Since we have created **the book**, why stop at recording monetary transactions?
- We can encode in the book's pages **arbitrary relations** between persons.
- Furthermore, scribes, can perform tasks such as verifying that stakeholders **comply** to contractual obligations ... and **take action** if they do not.

Smart Contract



Smart Contract Operation

- A smart contract is a **piece of code** written in a formal language that records all terms for a certain engagement between a set of persons, “stakeholders.”
- Stakeholders are identified by their **accounts**.
- The smart contract has a **public state**.
- The smart contract **self executes** each time a certain trigger condition is fulfilled.

Questions to Consider

- How are pages created? since the book is empty at the beginning, where do the money come from?

proofs of work

- How is it possible to sign something digitally?

digital signatures

- How does a page properly refer to the previous page?

hash functions

Hash Functions

Hash Functions

- An algorithm that produces a fingerprint of a file.

$$\mathcal{H} : \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$$

- what are the required properties (traditionally):
Efficiency.
A good spread for various input distributions.
- What are Security/Cryptographic considerations?

Collisions

- Collision Attack: Find $x, y : \mathcal{H}(x) = \mathcal{H}(y)$
- Second pre-image Attack: Find $y : \mathcal{H}(x) = \mathcal{H}(y)$

For given x

Birthday Paradox

- How many people should be in a room so that the probability that two of them share a birthday becomes larger than 50% ?

Paradox Explained

n possible dates

k people

$$\Pr[\neg Col] =$$

$$\frac{n}{n} \frac{n-1}{n} \frac{n-2}{n} \dots \frac{n-k+1}{n} = \prod_{\ell=1}^k \left(1 - \frac{\ell}{n}\right)$$
$$\leq \exp\left(-\frac{1}{n} \sum_{\ell=1}^k \ell\right) = \exp(-k(k+1)/2n)$$

$$\Pr[Col] = \frac{1}{2} \Rightarrow k \approx 1.177\sqrt{n}$$

Finding Collisions via the BP

- Randomly sample pairs of the form $\langle x, \mathcal{H}(x) \rangle$
- Store in table and sort according to 2nd coordinate.
- Perform linear pass to see whether there are elements with an equal 2nd coordinate

Analysis

- For k elements.
- Running time is $O(k \log k)$
- Choose k as $1.177\sqrt{2^\lambda}$

Pre-Image Attack

- Given $\mathcal{H}(m)$ $m \in \{0, 1\}^t$

Find an element of $\mathcal{H}^{-1}(\mathcal{H}(m))$

- Generic algorithms tries all possible values. How many?

$$O(2^t)$$

Constructing Hash Functions

- Relates to the notion of one-way functions.

$$f : X \rightarrow Y$$

easy : given x find $f(x)$

hard : given $f(x)$ sample $f^{-1}(f(x))$

Word of Caution

Existence of one-way functions
implies that :

$$P \neq NP$$

The single most important open question
in Computer Science today

Hash Function Implementations


- **Retired.** MD5, SHA1.
- **Current.** SHA2, SHA3, available for 224,256,384,512 bits fingerprints.
- **Bitcoin.** Uses SHA2 with 256 bits output, SHA-256.

Digital Signatures

What is a Signature?

- Can be produced by one specified entity.
- Can be verified by anyone (that is suitably “equipped” and “initialized”).
- Cannot be forged on a new message even if multiple signatures have been transmitted.

Pen and Paper Signatures



What are the assumptions when you verify them?
When you produce them ?

Digital Signatures

Three algorithms (**KeyGen**, **Sign**, **Verify**)

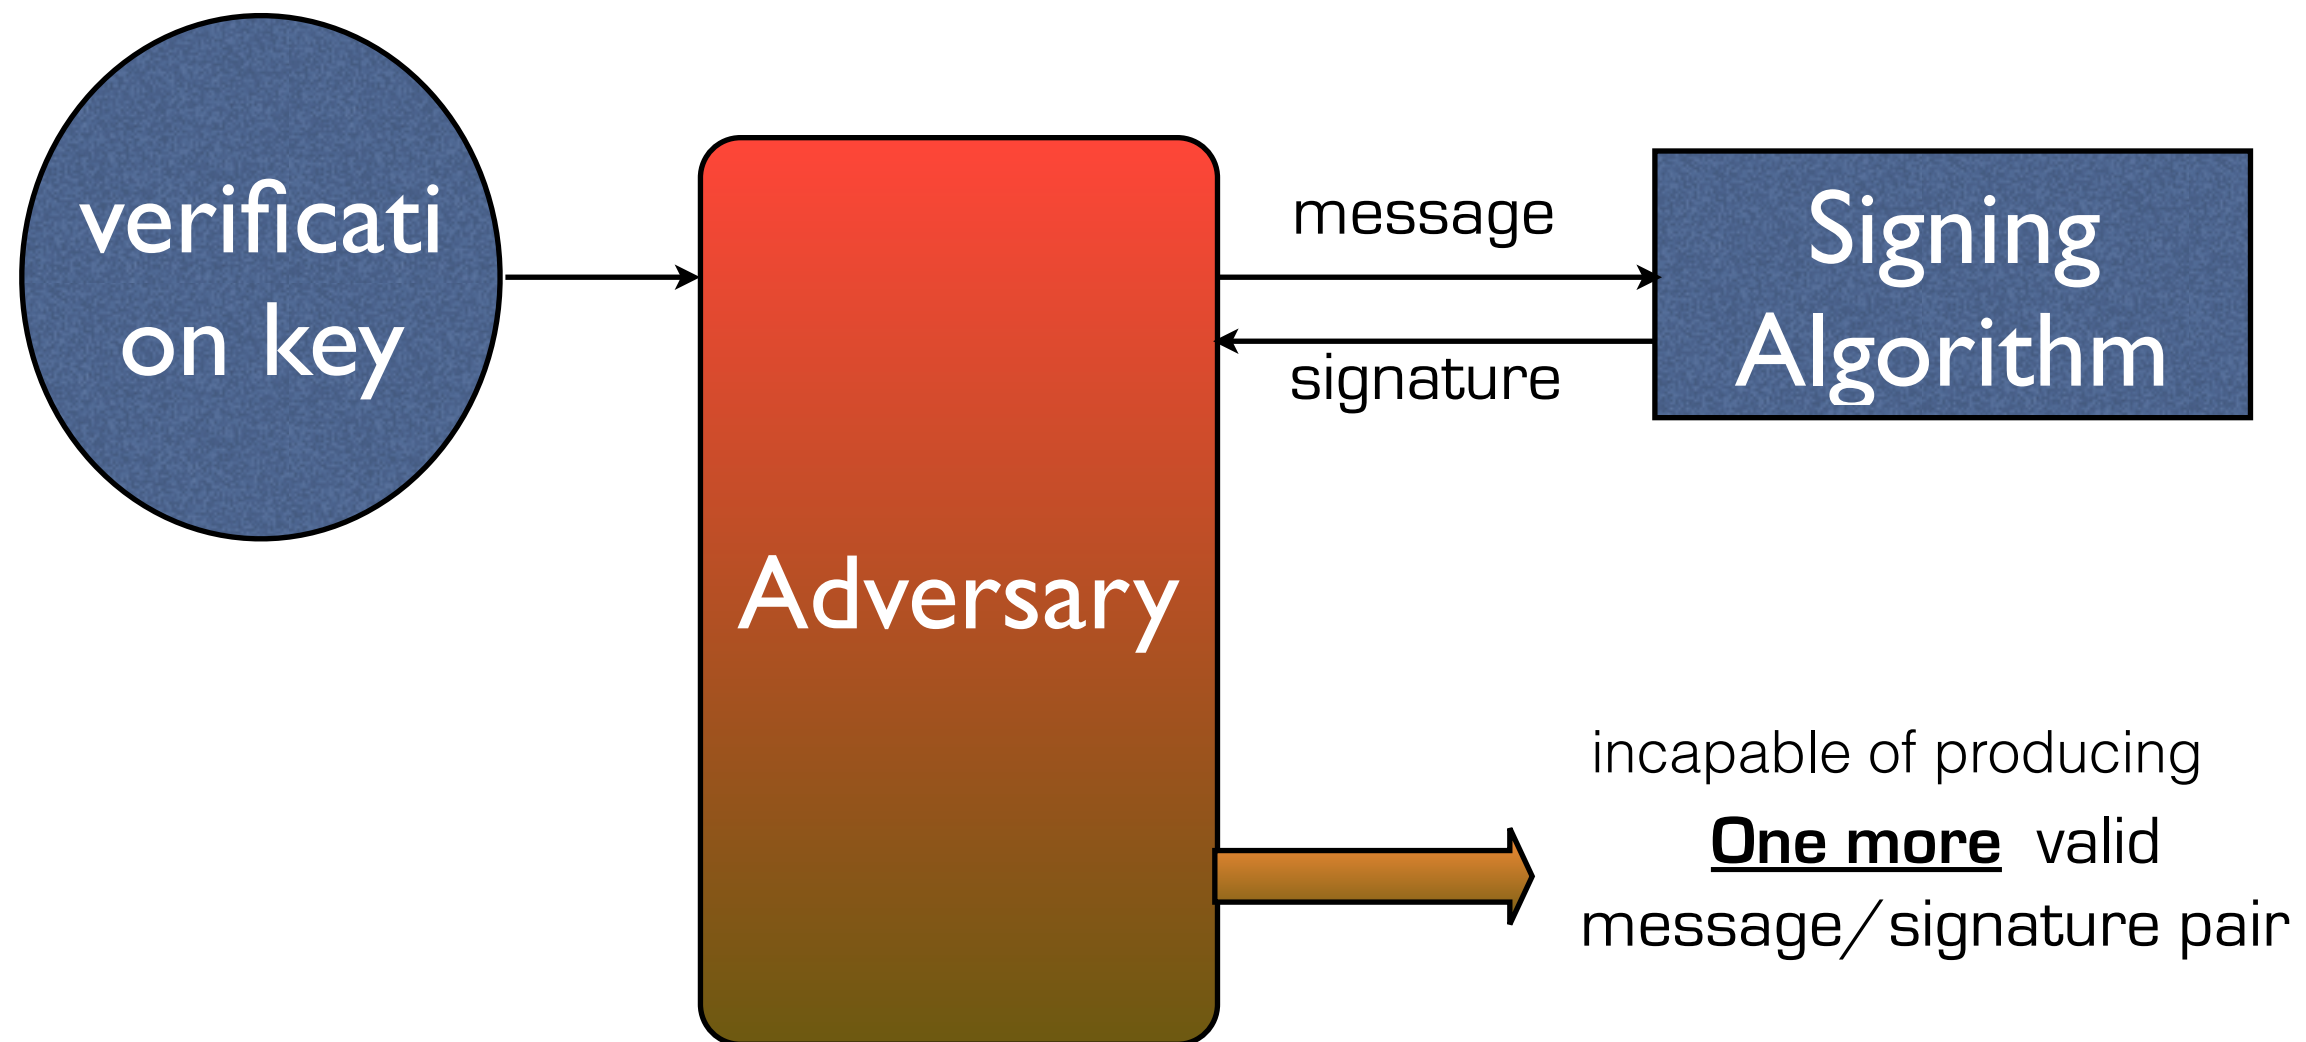
KeyGen : takes as input the *security parameter*.
returns the signing-key and verification-key.

Sign : takes as input the *signing-key* and
the *message* to be signed and
returns a signature.

Verify : takes as input the *verification-key*,
a *message* and a *signature* on the message and
returns either True or False.

Digital Signature Security

Existential Unforgeability under a Chosen Message Attack (EU-CMA)



Constructing Digital Signatures

- Major challenge:
 - what prevents the adversary from learning how to *sign* messages by analyzing the *verification-key*?

Example Construction : using hash functions

One-Time Signatures

$$\mathcal{H} : \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$$

KeyGen

$$sk = \langle s_{i,b} \rangle_{i \in \{1, \dots, \lambda\}, b \in \{0, 1\}} \quad vk = \langle \mathcal{H}(s_{i,b}) \rangle_{i \in \{1, \dots, \lambda\}, b \in \{0, 1\}}$$

Sign

$$\mathcal{H}(m) = w = w_1 \dots w_\lambda$$

$$\sigma = \langle s_{1,w_1}, \dots, s_{\lambda,w_\lambda} \rangle$$

Verify Hash strings from σ and compare to the hashes stored in the vk at the positions dictated by $\mathcal{H}(m)$

Digital Signature Constructions

- Based on the RSA (Rivest Shamir Adleman), one way trapdoor function (with hardness that relates to the factoring problem).
 - The RSA algorithm
- Based on the discrete-logarithm problem.
 - the DSA algorithm
- **Bitcoin.** Uses ECDSA, a DSA variant over elliptic curve groups.

bitcoin transactions

- Typical transaction
 - input: contains a signature and public-key
 - output: contains a verification procedure

General Format

general format of a Bitcoin transaction (inside a block)

Field	Description	Size
Version no	currently 1	4 bytes
In-counter	positive integer $VI = \text{VarInt}$	1 - 9 bytes
list of inputs	the first input of the first transaction is also called "coinbase" (its content was ignored in earlier versions)	<in-counter>-many inputs
Out-counter	positive integer $VI = \text{VarInt}$	1 - 9 bytes
list of outputs	the outputs of the first transaction spend the mined bitcoins for the block	<out-counter>-many outputs
lock_time	if non-zero and sequence numbers are $< 0xFFFFFFFF$: block height or timestamp when transaction is final	4 bytes

Input and Output Scripts

Input:

Previous tx: f5d8ee39a430901c91a5917b9f2dc19d6d1a0e9cea205b009ca73dd04470b9a6

Index: 0

scriptSig:

304502206e21798a42fae0e854281abd38bacd1aeed3ee3738d9e1446618c4571d10

90db022100e2ac980643b0b82c0e88ffdfec6b64e3e6ba35e7ba5fdd7d5d6cc8d25c6b241501

Output:

Value: 5000000000

scriptPubKey: OP_DUP OP_HASH160 404371705fa9bd789a2fcd52d2c580b65d35549d

OP_EQUALVERIFY OP_CHECKSIG

The input imports 50 BTC from output #0 of tx f5d..
and sends them to a Bitcoin address 404...

Verifying a transaction

```
scriptPubKey: OP_DUP OP_HASH160 <pubKeyHash>  
              OP_EQUALVERIFY OP_CHECKSIG
```

```
scriptSig: <sig> <pubKey>
```

Transaction Processing

- Uses a **stack** data structure:
 - items can be pushed to the stack and popped from the stack.
 - following LIFO (“last in first out”) order.

Transaction processing

Stack	Script	Description
Empty.	<sig> <pubKey> OP_DUP OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG	scriptSig and scriptPubKey are combined.
<sig> <pubKey>	OP_DUP OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG	Constants are added to the stack.
<sig> <pubKey> <pubKey>	OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG	Top stack item is duplicated.
<sig> <pubKey> <pubHashA>	<pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG	Top stack item is hashed.
<sig> <pubKey> <pubHashA> <pubKeyHash>	OP_EQUALVERIFY OP_CHECKSIG	Constant added.
<sig> <pubKey>	OP_CHECKSIG	Equality is checked between the top two stack items.
true	Empty.	Signature is checked for top two stack items.

More general transactions

- Transactions can include arbitrary script operations.

e.g., pay to script hash

```
scriptPubKey: OP_HASH160 <scriptHash> OP_EQUAL  
scriptSig: ..signatures... <serialized script>
```

(instead of):

```
scriptPubKey: OP_DUP OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG  
scriptSig: <sig> <pubKey>
```

script operations

- Some operations that are allowed:
 - basic 32-bit arithmetic (addition subtraction)
 - computation of SHA256[.]
 - verify equality
 - calculate length of string
 - Verify signature (ECDSA)

End of Lecture 01

- Next time :
 - The consensus layer.
 - Basic Properties.
 - Proof of work.