

# Blockchains and Distributed Ledgers Lecture 08

Aggelos Kiayias



THE UNIVERSITY  
*of* EDINBURGH

# Lecture 08

- Permissionless vs. Permissioned Ledgers.
- BFT and PoS-based approaches for permissioned ledgers.
- Dynamic participation.

# Permissionless Protocols

- Bitcoin and similar PoW-based blockchain protocols provide a **permissionless** setting:
  - Anyone can participate in the protocol and receive BTC as rewards by performing the PoW-based mining operation.
  - The mechanism of **pouring** currency in the system via proof of work, makes it feasible for anyone (possessing sufficient hashing power) to participate.
  - The ledger itself is public, readable and writeable by anyone (the latter assuming one **possesses** bitcoin)

# Permissioned Protocols

- Participation is restricted:
  - Producing transactions and/or blocks can only be performed after being authorized by the other nodes.
- In their simplest form the set of nodes is **static**: the set of nodes implementing the protocol is fixed and determined at the onset of protocol execution.

# Permissioning How-To

- Most straight approach: employ a PKI (=public-key infrastructure).
- Based on digital signatures / authentication protocols.
- **Certificate authorities** can authorize other entities.
  - authorization includes a signature from the CA on the entity's public-key, identity information etc.
- Sharing certificate authority information is necessary; (how, where?)

# X.509 Certificates

- Internet standard since 1988.
- Hierarchical.
- <http://www.ietf.org/rfc/rfc3280.txt>

# Structure of X.509 Certificates

Version
Serial Number
Algorithm / Parameters
Issuer
Period of Validity: not before date not after date
Subject
Algorithm/ Parameters/ Key
x509v3 extensions
...
<i>Signature</i>

**X.509**  
does not  
specify  
cryptographic  
algorithms

# Example, I

Certificate:

Data:

Version: 3 (0x2)

Serial Number:

13:86:35:4d:1d:3f:06:f2:c1:f9:65:05:d5:90:1c:62

Signature Algorithm: sha1WithRSAEncryption

Issuer: C=US, O=VISA, OU=Visa International Service Association, CN=Visa eCommerce Root

Validity

Not Before: Jun 26 02:18:36 2002 GMT

Not After : Jun 24 00:16:12 2022 GMT

Subject: C=US, O=VISA, OU=Visa International Service Association, CN=Visa eCommerce Root

Subject Public Key Info:

Public Key Algorithm: rsaEncryption

RSA Public Key: (2048 bit)

Modulus (2048 bit):

00:af:57:de:56:1e:6e:a1:da:60:b1:94:27:cb:17:  
db:07:3f:80:85:4f:c8:9c:b6:d0:f4:6f:4f:cf:99:  
d8:e1:db:c2:48:5c:3a:ac:39:33:c7:1f:6a:8b:26:  
3d:2b:35:f5:48:b1:91:c1:02:4e:04:96:91:7b:b0:  
33:f0:b1:14:4e:11:6f:b5:40:af:1b:45:a5:4a:ef:  
7e:b6:ac:f2:a0:1f:58:3f:12:46:60:3c:8d:a1:e0:  
7d:cf:57:3e:33:1e:fb:47:f1:aa:15:97:07:55:66:  
a5:b5:2d:2e:d8:80:59:b2:a7:0d:b7:46:ec:21:63:

 **Hierarchical name spec**

 **The RSA Algorithm is used;  
the DSA is also an option**

**This is a self-signed  
root certificate**



# Example, II

note:

**public-exponent,  
corresponding to RSA  
function: the public-key  
is a large integer N  
that is a product of two  
primes and a small  
exponent.**

4a:ea:db:df:72:38:8c:f3:96:bd:f1:17:bc:d2:ba:  
3b:45:5a:c6:a7:f6:c6:17:8b:01:9d:fc:19:a8:2a:  
83:16:b8:3a:48:fe:4e:3e:a0:ab:06:19:e9:53:f3:  
80:13:07:ed:2d:bf:3f:0a:3c:55:20:39:2c:2c:00:  
69:74:95:4a:bc:20:b2:a9:79:e5:18:89:91:a8:dc:  
1c:4d:ef:bb:7e:37:0b:5d:fe:39:a5:88:52:8c:00:  
6c:ec:18:7c:41:bd:f6:8b:75:77:ba:60:9d:84:e7:  
fe:2d

Exponent: 65537 (0x10001)

X509v3 extensions:

X509v3 Basic Constraints: critical  
CA:TRUE

X509v3 Key Usage: critical  
Certificate Sign, CRL Sign

X509v3 Subject Key Identifier:

15:38:83:0F:3F:2C:3F:70:33:1E:CD:46:FE:07:8C:20:E0:D7:C3:B7

Signature Algorithm: sha1WithRSAEncryption

5f:f1:41:7d:7c:5c:08:b9:2b:e0:d5:92:47:fa:67:5c:a5:13:  
c3:03:21:9b:2b:4c:89:46:cf:59:4d:c9:fe:a5:40:b6:63:cd:

.....



Note that SHA-1 is broken now : <https://shattered.io>

# Certificate Encoding

- Certificates are binary files.
- They can be Base64 encoded to be stored/transported as text files.
- This encoding is called PEM from (privacy enhanced mail)

-----BEGIN CERTIFICATE-----

```
MIIFXzCCBEegAwIBAgIHK5Bmkq9qPzANBgkqhkiG9w0BAQUFADCByjELMAkGA1UE
BhMCVVMxEDA0BgNVBAGTB0FyaXpvcnkvZ29tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIERhZGR5
BgNVBAoTEUdvRGFkZHZhY29tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIERhZGR5
aWNhdGVzLmdvZGFkZHZhY29tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIERhZGR5
IFNlY3VyZSBDZXJ0aWZpY2F0aW9uIEF1dGhvcml0eTERMA8GA1UEBRMIMDc5Njky
ODcwHhcNMTEwMTkyMzIyWhcNMTEwMTkyMzIyMTMyMTEyWjBZMRkwFwYDVQKDBAq
LmVuZ3IudWNvb4uZWRLMSEwHwYDVQQLDBhEb21haW4gQ29udHJvbCBWYXpZGF0
ZWQxGTAXBgNVBAMMEC0uZW5nci5lY29ubi5lZHUwggEiMA0GCSqGSIb3DQEBAQUA
A4IBDwAwggEKAoIBAQC0gAbFE4T0s2KA/Wm9rZ5Q+3G6teSFdjmxZynMeCch+/e/
bSudinQ0oJShqXkyiwrGalbrr3pbQkGLn7AUdli+479IM0q2tFpvZ9SgCVsp2Jk1
/6MqBleBTKfVGfVEz6+YzkwcoiEe7cY+deOeIG1ZFVlnI3kw19V42gk+ZU25oVgL
uiJB2RRD3djpaVltxMVyCnzHa0u2eEHhmh87faIA6EckEW/cWroVuBZF5Oojsr5c
iXgLIwImQn+Cmo7NQVGv4pKqJVnHPyFX8w+CcLJ8GonVxQeKleh02DycpDoi2SpP
/hAZ1Vh9mEfftpCikXPwFlunoZAmn9mWUtpcyMnAgMBAAGjggG4MIIBtDAPBgNV
HRMBAf8EBTADAQEAMBA0GA1UdJQQWMBQGCCsGAQUFBwMBBggrBgEFBQcDAjA0BgNV
HQ8BAf8EBAMCBaAwMwYDVR0fBCwwKjAooCagJIYiaHR0cDovL2Nybc5nb2RhZGR5
LmNvbS9nZHMxLTI4LmNybDBNBgNVHSAERjBEMEIGC2CGSAGG/W0BBxcBMDMwMQYI
KwYBBQUHAgEJWh0dHBzOi8vY2VydHMuZ29kYWRkeS5jb20vcvVwb3NpdG9yeS8w
gYAGCCsGAQUFBwEBBHQwcjAkBggrBgEFBQcwAYYYaHR0cDovL29jc3AuZ29kYWRk
eS5jb20vMEoGCCsGAQUFBzACHj5odHRwOi8vY2VydGlmawWNhdGVzLmdvZGFkZHZhY2
9tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIERhZGR5IFNlY3VyZSBDZXJ0aWZpY2
F0aW9uIEF1dGhvcml0eTERMA8GA1UEBRMIMDc5NjkyODcwHhcNMTEwMTkyMzIyWhc
NMTEwMTkyMzIyMTMyMTEyWjBZMRkwFwYDVQKDBAqLmVuZ3IudWNvb4uZWRLMSEwH
wYDVQQLDBhEb21haW4gQ29udHJvbCBWYXpZGF0ZWQxGTAXBgNVBAMMEC0uZW5nci5
lY29ubi5lZHUwggEiMA0GCSqGSIb3DQEBAQUAA4IBDwAwggEKAoIBAQC0gAbFE4T0s
2KA/Wm9rZ5Q+3G6teSFdjmxZynMeCch+/e/bSudinQ0oJShqXkyiwrGalbrr3pbQkGL
n7AUdli+479IM0q2tFpvZ9SgCVsp2Jk1/6MqBleBTKfVGfVEz6+YzkwcoiEe7cY+de
OeIG1ZFVlnI3kw19V42gk+ZU25oVgLuiJB2RRD3djpaVltxMVyCnzHa0u2eEHhmh87
faIA6EckEW/cWroVuBZF5Oojsr5ciXgLIwImQn+Cmo7NQVGv4pKqJVnHPyFX8w+CcLJ
8GonVxQeKleh02DycpDoi2SpP/hAZ1Vh9mEfftpCikXPwFlunoZAmn9mWUtpcyMnAgM
BAAGjggG4MIIBtDAPBgNVHRMBAf8EBTADAQEAMBA0GA1UdJQQWMBQGCCsGAQUFBwMB
BggrBgEFBQcDAjA0BgNVHQ8BAf8EBAMCBaAwMwYDVR0fBCwwKjAooCagJIYiaHR0cD
ovL2Nybc5nb2RhZGR5LmNvbS9nZHMxLTI4LmNybDBNBgNVHSAERjBEMEIGC2CGSAGG
/W0BBxcBMDMwMQYIKwYBBQUHAgEJWh0dHBzOi8vY2VydHMuZ29kYWRkeS5jb20vcv
Vwb3NpdG9yeS8wYAGCCsGAQUFBwEBBHQwcjAkBggrBgEFBQcwAYYYaHR0cDovL29jc3
AuZ29kYWRkeS5jb20vMEoGCCsGAQUFBzACHj5odHRwOi8vY2VydGlmawWNhdGVzLmd
vZGFkZHZhY29tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIERhZGR5IFNlY3VyZSBD
ZXJ0aWZpY2F0aW9uIEF1dGhvcml0eTERMA8GA1UEBRMIMDc5NjkyODcwHhcNMTEw
MTkyMzIyWhcNMTEwMTkyMzIyMTMyMTEyWjBZMRkwFwYDVQKDBAqLmVuZ3IudWNvb4
uZWRLMSEwHwYDVQQLDBhEb21haW4gQ29udHJvbCBWYXpZGF0ZWQxGTAXBgNVBAMME
C0uZW5nci5lY29ubi5lZHUwggEiMA0GCSqGSIb3DQEBAQUAA4IBDwAwggEKAoIBAQC
0gAbFE4T0s2KA/Wm9rZ5Q+3G6teSFdjmxZynMeCch+/e/bSudinQ0oJShqXkyiwrGa
lbrr3pbQkGLn7AUdli+479IM0q2tFpvZ9SgCVsp2Jk1/6MqBleBTKfVGfVEz6+Yzkw
coiEe7cY+deOeIG1ZFVlnI3kw19V42gk+ZU25oVgLuiJB2RRD3djpaVltxMVyCnzHa
0u2eEHhmh87faIA6EckEW/cWroVuBZF5Oojsr5ciXgLIwImQn+Cmo7NQVGv4pKqJVn
HPyFX8w+CcLJ8GonVxQeKleh02DycpDoi2SpP/hAZ1Vh9mEfftpCikXPwFlunoZAmn
9mWUtpcyMnAgMBAAGjggG4MIIBtDAPBgNVHRMBAf8EBTADAQEAMBA0GA1UdJQQWMB
QGCCsGAQUFBwMBBggrBgEFBQcDAjA0BgNVHQ8BAf8EBAMCBaAwMwYDVR0fBCwwKjA
ooCagJIYiaHR0cDovL2Nybc5nb2RhZGR5LmNvbS9nZHMxLTI4LmNybDBNBgNVHSAE
RjBEMEIGC2CGSAGG/W0BBxcBMDMwMQYIKwYBBQUHAgEJWh0dHBzOi8vY2VydHMuZ2
9kYWRkeS5jb20vcvVwb3NpdG9yeS8wYAGCCsGAQUFBwEBBHQwcjAkBggrBgEFBQcw
AYYYaHR0cDovL29jc3AuZ29kYWRkeS5jb20vMEoGCCsGAQUFBzACHj5odHRwOi8vY2
VydGlmawWNhdGVzLmdvZGFkZHZhY29tL3JlcG9zaXRvcnkxMDAuBgNVBAMTJ0dvIER
hZGR5IFNlY3VyZSBDZXJ0aWZpY2F0aW9uIEF1dGhvcml0eTERMA8GA1UEBRMIMDc5
NjkyODcwHhcNMTEwMTkyMzIyWhcNMTEwMTkyMzIyMTMyMTEyWjBZMRkwFwYDVQKDB
AqLmVuZ3IudWNvb4uZWRLMSEwHwYDVQQLDBhEb21haW4gQ29udHJvbCBWYXpZGF0Z
WQxGTAXBgNVBAMMEC0uZW5nci5lY29ubi5lZHUwggEiMA0GCSqGSIb3DQEBAQUA=
```

-----END CERTIFICATE-----

# Binary Encoding

- The elements of certificate are encoded using DER (distinguished encoding rules)

Items are stored in TLV format : triplets <Type,Length,Value>

```
1. 30 23      ; SEQUENCE (23 Bytes)
2. | | 31 0f    ; SET (f Bytes)
3. | | | 30 0d   ; SEQUENCE (d Bytes)
4. | | | 06 03   ; OBJECT_ID (3 Bytes)
5. | | | | 55 04 03
6. | | | |      ; 2.5.4.3 Common Name (CN)
7. | | | 13 06   ; PRINTABLE_STRING (6 Bytes)
8. | | | 54 65 73 74 43 4e      ; TestCN
9. | | |      ; "TestCN"
10. | | 31 10    ; SET (10 Bytes)
11. | | 30 0e    ; SEQUENCE (e Bytes)
12. | | 06 03   ; OBJECT_ID (3 Bytes)
13. | | | 55 04 0a
14. | | |      ; 2.5.4.10 Organization (O)
15. | | 13 07   ; PRINTABLE_STRING (7 Bytes)
16. | | 54 65 73 74 4f 72 67      ; TestOrg
17. | |      ; "TestOrg"
```

# Digital Signatures and Certificates

- A certificate contains a digital signature.
- Recall that cryptographic design of digital signatures involves typically:
  - A cryptographic signing operation that acts on a fixed input of a specific type and has a public-verifiability feature.
  - A cryptographic hash function that takes arbitrary strings and maps them to the data type suitable for the signing operation.
- Common setting today: SHA2 with RSA/DSA.

# Certificate Considerations

- All computer systems come with preloaded certificates from certificate authorities. This provide a **setup assumption**.
- Certificates need to be revoked in case the corresponding secret keys become exposed or the algorithms used are not safe anymore.
- In a blockchain system, certificate information can be provided as part of the genesis block.

# Secure Channels and Certificates

- Possession of mutually acceptable certificates not only permits authenticated communication (exchanging signed mechanism between two entities) but also allows building a secure channel
- Protocol **TLS 1.2** is used to build such secure channel.
- It relies on cryptographic protocols such as the Diffie Hellman key exchange. It can ensure the confidentiality of the data exchanged.

# Static Permissioned Blockchain

- All participants are identified by self-signed certificates in the genesis block.
- The set of participants remains the same throughout the execution.
- This is the simplest form of a PKI / public-key directory.

# Recall: Robust Transaction Ledger

- persistence: Transactions are organized in a “log” and honest nodes agree on it.
- liveness: New transactions are included in the log nodes, after a suitable period of time.



# In more detail

- Common prefix / persistence / consistency of the LOG:
  - For each two nodes that maintains a log, and at any two times  $t_1 \leq t_2$ , it holds that LOG1 is a prefix of LOG2  
(where LOG2 includes pending transactions)
- Liveness of the log:
  - An honestly generated transaction will become part of all honest parties' LOGs after a specified time window  $u$ .

# Permissioning

- Prior to system operation the nodes register their certificates that are included in the genesis block.
- Using such certificates, all the nodes are capable of authenticating each participant and allowing interaction with the LOG in a way that is prescribed by the participants' credentials.

# Centralized Permissioned Ledger

- One of participants acts as a server and maintains the LOG.
- Readers and writers to the LOG authenticate with the server and can perform read and write operations.
- Consistency of the LOG is guaranteed assuming the server is trusted.
- Liveness of the LOG is guaranteed assuming the server is trusted and functional.

# Bitcoin's Permissionless Setting

- The genesis block contains no certificate information.
- Reading from the LOG is open (anyone can do it without credentials).
- Writing to the LOG can only be done in specific ways (issuing transactions).
  - Nodes can obtain valid credentials (accounts) by generating a public and secret-key and either mining a block (which will reward their account with BTC) or buy BTC from another node.
  - Once the LOG records their account credit, they can issue transactions.
  - In essence: crediting a bitcoin account is creating a certificate that imparts the account holder with certain permissions w.r.t. the ledger.

# Distributed Permissioned Ledger, I

- A number of servers maintain the ledger LOG individually.
- Each share the same genesis block that identifies all participants.
- Assuming a synchronous operation, at each round, Readers and Writers authenticate with the servers and interact with the LOG in a prescribed fashion.

# Distributed Permissioned Ledger, II

- Readers authenticate to each server and obtain Read access.
- Writers authenticate to each server and provide their inputs.
- Servers run a **consensus protocol** to agree what inputs should be included in the LOG.

# Reader/Writer Management

- Readers and Writers can authenticate to each server referring to the information in the genesis block.
- It is possible to introduce additional readers and writers by suitably issuing certificates to other users.
- Note that each participant would then need to show a valid certificate chain that establishes her privileges for the specific read or write access that is requested.

# Read requests

- Is it possible to restrict read requests as in the centralized setting?
- Nodes can maintain blocks of transactions private and issue them only to users that are authenticated.
- The TLS protocol can be used to build a secure channel between the reader and the responding node.
- Note that the above would require that all servers remain honest (as they all share the LOG).



# A classical BFT-based approach for permissioned ledgers

- Focus on write requests next. We want to ensure LOG liveness and consistency.
- We will apply a “byzantine fault tolerant” (BFT) agreement protocol that uses two important tools:
  - a graded broadcast.
  - and a common coin binary consensus protocol.

# Graded Broadcast

- Parties involved : a single sender and several receivers.
- The  $i$ -th receiver outputs  $(M_i, G_i)$ .
- The value  $G_i$  is in  $\{0,1,2\}$ .
  - If the sender is honest then  $M_i=M_j$  for all  $i,j$  and  $G_i=2$ .
  - If the sender is malicious and one receiver outputs  $G_i=2$  then other honest receivers output  $G_i \geq 1$ .

# Graded Broadcast Construction, I

- **Round 1.** The sender sends the message  $M$  to all receivers.
- **Round 2.** The  $i$ -th receiver obtains  $M1_i$  from round 1 and sends it to all receivers.
- **Round 3.** The  $i$ -th receiver obtains  $M2_{ji}$  from the  $j$ -th receiver in round 2 and performs the following:  
if there is a single message that was sent by  $2n/3$  receivers then send it to all receivers. Else send nothing.

# Graded Broadcast Construction, II

- The  $i$ -th receiver obtains  $M_{3ji}$  from the  $j$ -th receiver in round 3.
- If there is a single message that was sent by more than  $2n/3$  receivers output that message as  $M_i$  and set  $G_i=2$ .
- If there is a single message that was sent by more than  $n/3$  receivers output that message as  $M_i$  and set  $G_i=1$ .
- In any other case output *fail* and  $G_i=0$ .

# Graded Broadcast Construction, III

- Analysis. Assume that malicious parties are  $t < n/3$ .
- If the sender is honest, then each receiver will receive the same message  $\geq 2n/3$  times in round 2 and 3. All honest receivers will output  $G_{i=2}$  and that message.

# Graded Broadcast IV

- **Lemma.** If two honest receivers send a message in round 3 it **must be** the *same*.
  - Indeed, if they send messages  $M \neq M'$ , they both have received them by at least  $2n/3$  receivers from round 2.
  - Given the above, observe that  $2n/3 - t > n/3$  honest parties have sent  $M$  in round 2.
  - Thus  $<n - n/3 = 2n/3$  parties are capable of sending  $M'$  (which is different than  $M$ ), leading to contradiction.

# Graded Broadcast, V

- Suppose the  $i$ -th receiver returns  $G_i=2$  and let  $M_i$  be the message it chooses. Consider the output  $(M_j, G_j)$ 
  - The  $i$ -th receiver has received the message from at least  $2n/3$  receivers in round 3.
  - $\Rightarrow$  More than  $n/3$  honest receivers have sent  $M$  in round 3. Thus the  $j$ -th receiver should produce a grade  $G_j \geq 1$ .
  - If there is another message  $M'$  sent by more  $n/3$  receivers in round 3, at least one of them is honest; this leads to a contradiction of the **lemma**. Thus  $M_j=M_i$ .

# From Graded Broadcast to a BFT-based ledger

- Execute  $n/3$  phases:
  - In each phase perform:
    - A designated sender organizes all valid transactions it collected as **M** and performs a graded broadcast.
    - A binary consensus protocol that determines whether everyone's grade is 2. If that is true each node signs the output to generate a public endorsement and appends **M** on their LOG (together with the signatures). Otherwise LOG remains the same.



# Recall Consensus Properties

- Termination  $\forall i \in \mathbf{H} (u_i \text{ is defined})$
- Agreement  $\forall i, j \in \mathbf{H} (u_i = u_j)$
- Validity  $\exists v (\forall i \in \mathbf{H} (v_i = v)) \implies (\forall i \in \mathbf{H} (u_i = v))$

# Common Coin Consensus, I

- *$B_i$  is the input of party  $i$  in  $\{0, 1\}$*
- **Round 1.** The  $i$ -th party sends  $B_i$  to all parties.
- **Round 2.** A randomly selected common coin  $c$  is made available.
- **Round 3.** The  $i$ -th party obtains  $B_i$  from round 0 and performs the following:
  - if 0 was sent by  $2n/3$  parties set  $B_i=0$ .
  - if 1 was sent by  $2n/3$  parties set  $B_i=1$ .
  - else set  $B_i=c$ .

# Common Coin Consensus, II

- **Validity.** Suppose all honest parties agree on their input initially; i.e., for all  $i, j$ ,  $B_i = B_j$ .
  - Then, given that  $t < n/3$ , they will set  $B$  to that bit in round 3.
- **Agreement.** Suppose that an honest party outputs 0 because  $2n/3$  sent 0 in round 1. It follows that  $> n/3$  honest parties had 0 as input prior to the protocol. As a result any other honest party will also output 0 with probability  $1/2$ . Similar argument can be made for 1.
- **Termination.** it will always happen in 3 rounds.

# Common Coin Consensus, III

- Repeat protocol for  $k$  times.
- After  $k$  repetitions, it is guaranteed that honest parties will reach agreement with probability  $1-2^{-k}$  (original input is only used in first rep)
- (note that they will not know they did).
- [more involved techniques exist that bring the number of repetitions to **expected constant** with detection of agreement].

# Producing a common coin

- Suppose that the common coin protocol is imperfect and serves a proper coin with probability  $\varepsilon$  and no other assumptions otherwise. E.g., with probability  $1-\varepsilon$ , the parties may produce a different adversarially selected value.
- Then, agreement in the consensus protocol will only hold with probability  $\varepsilon/2$ .
- Same analysis as before works but need to repeat protocol for  $k\varepsilon^{-1}$  times.
  - It is easy to implement common coin for small  $\varepsilon$ : have each party simply flip a coin. it holds that  $\varepsilon = 2^{-n+t+1}$ . [More involved techniques provide a large constant  $\varepsilon < 1$ ]

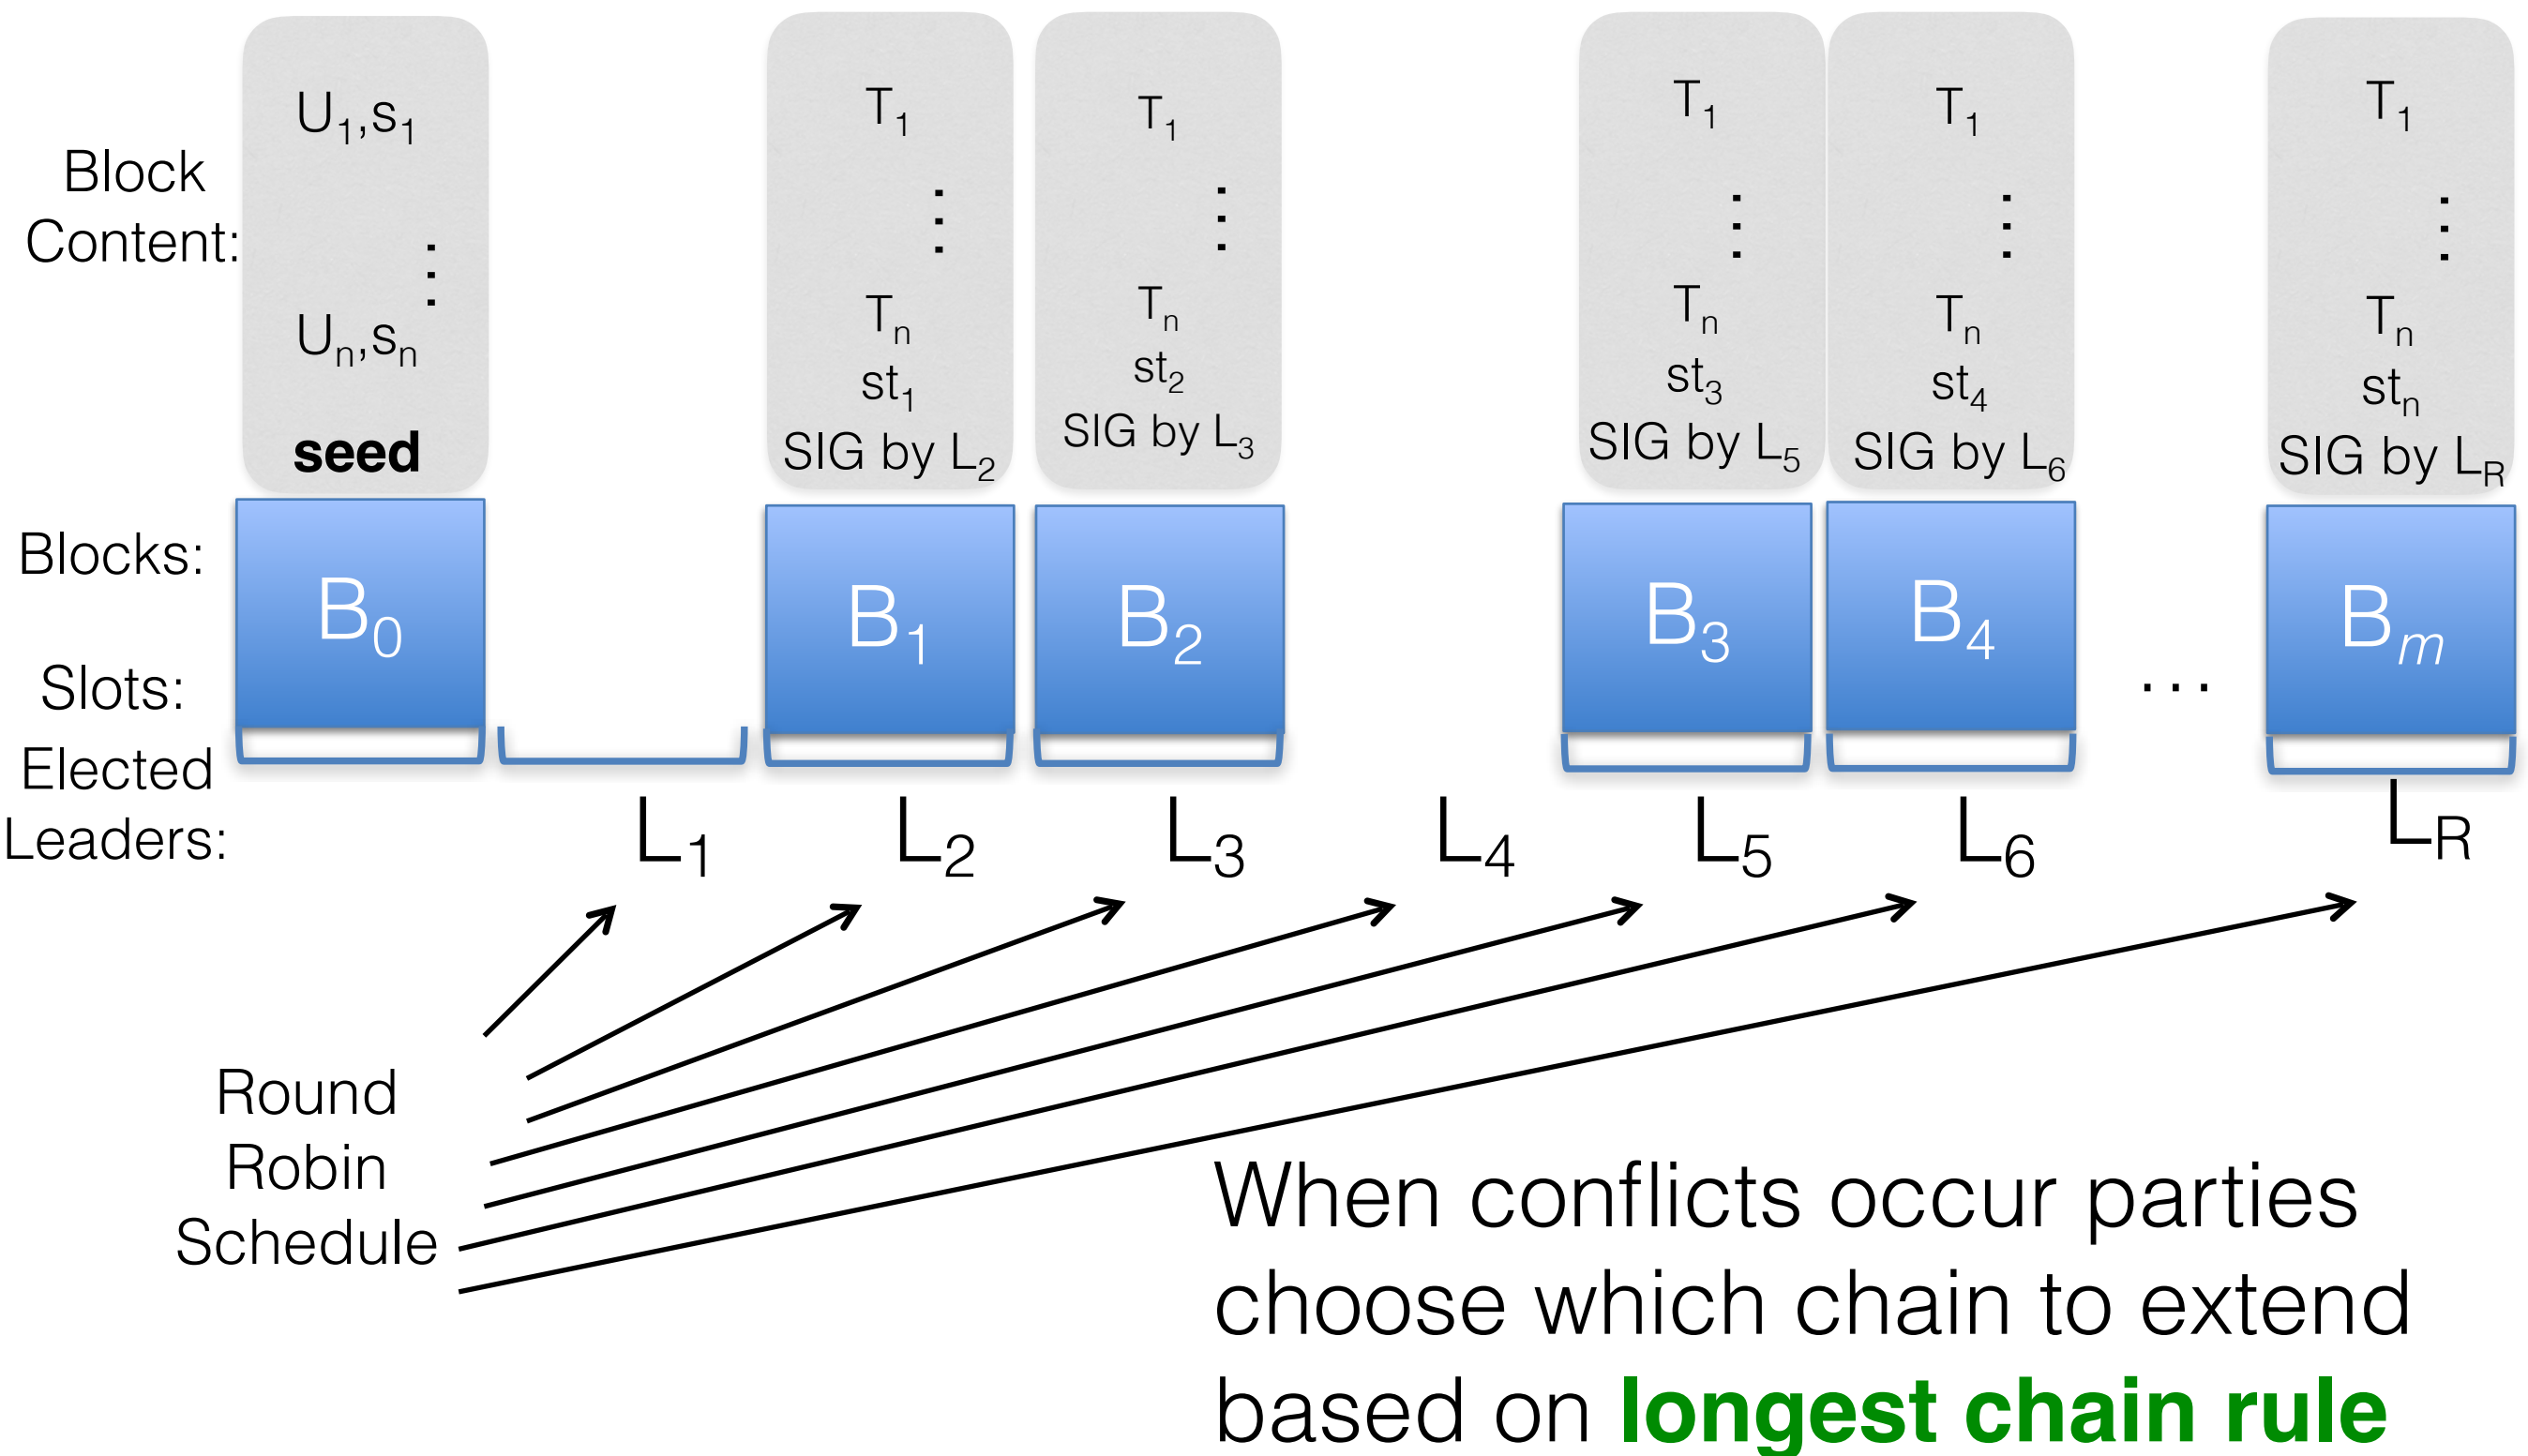
# Distributed Ledger

- Using the above protocol it is guaranteed that:
  - The LOG will grow by a new set of transactions.
  - The transactions are contributed by an honest server and thus include all new transactions received by that node.
  - The waiting time  $u$ , for a transaction to be included is proportional to the termination of the protocol that agrees on the next LOG entry; (exponential in our example but can be improved to constant).

# PoS-based Approach for Permissioned Ledgers

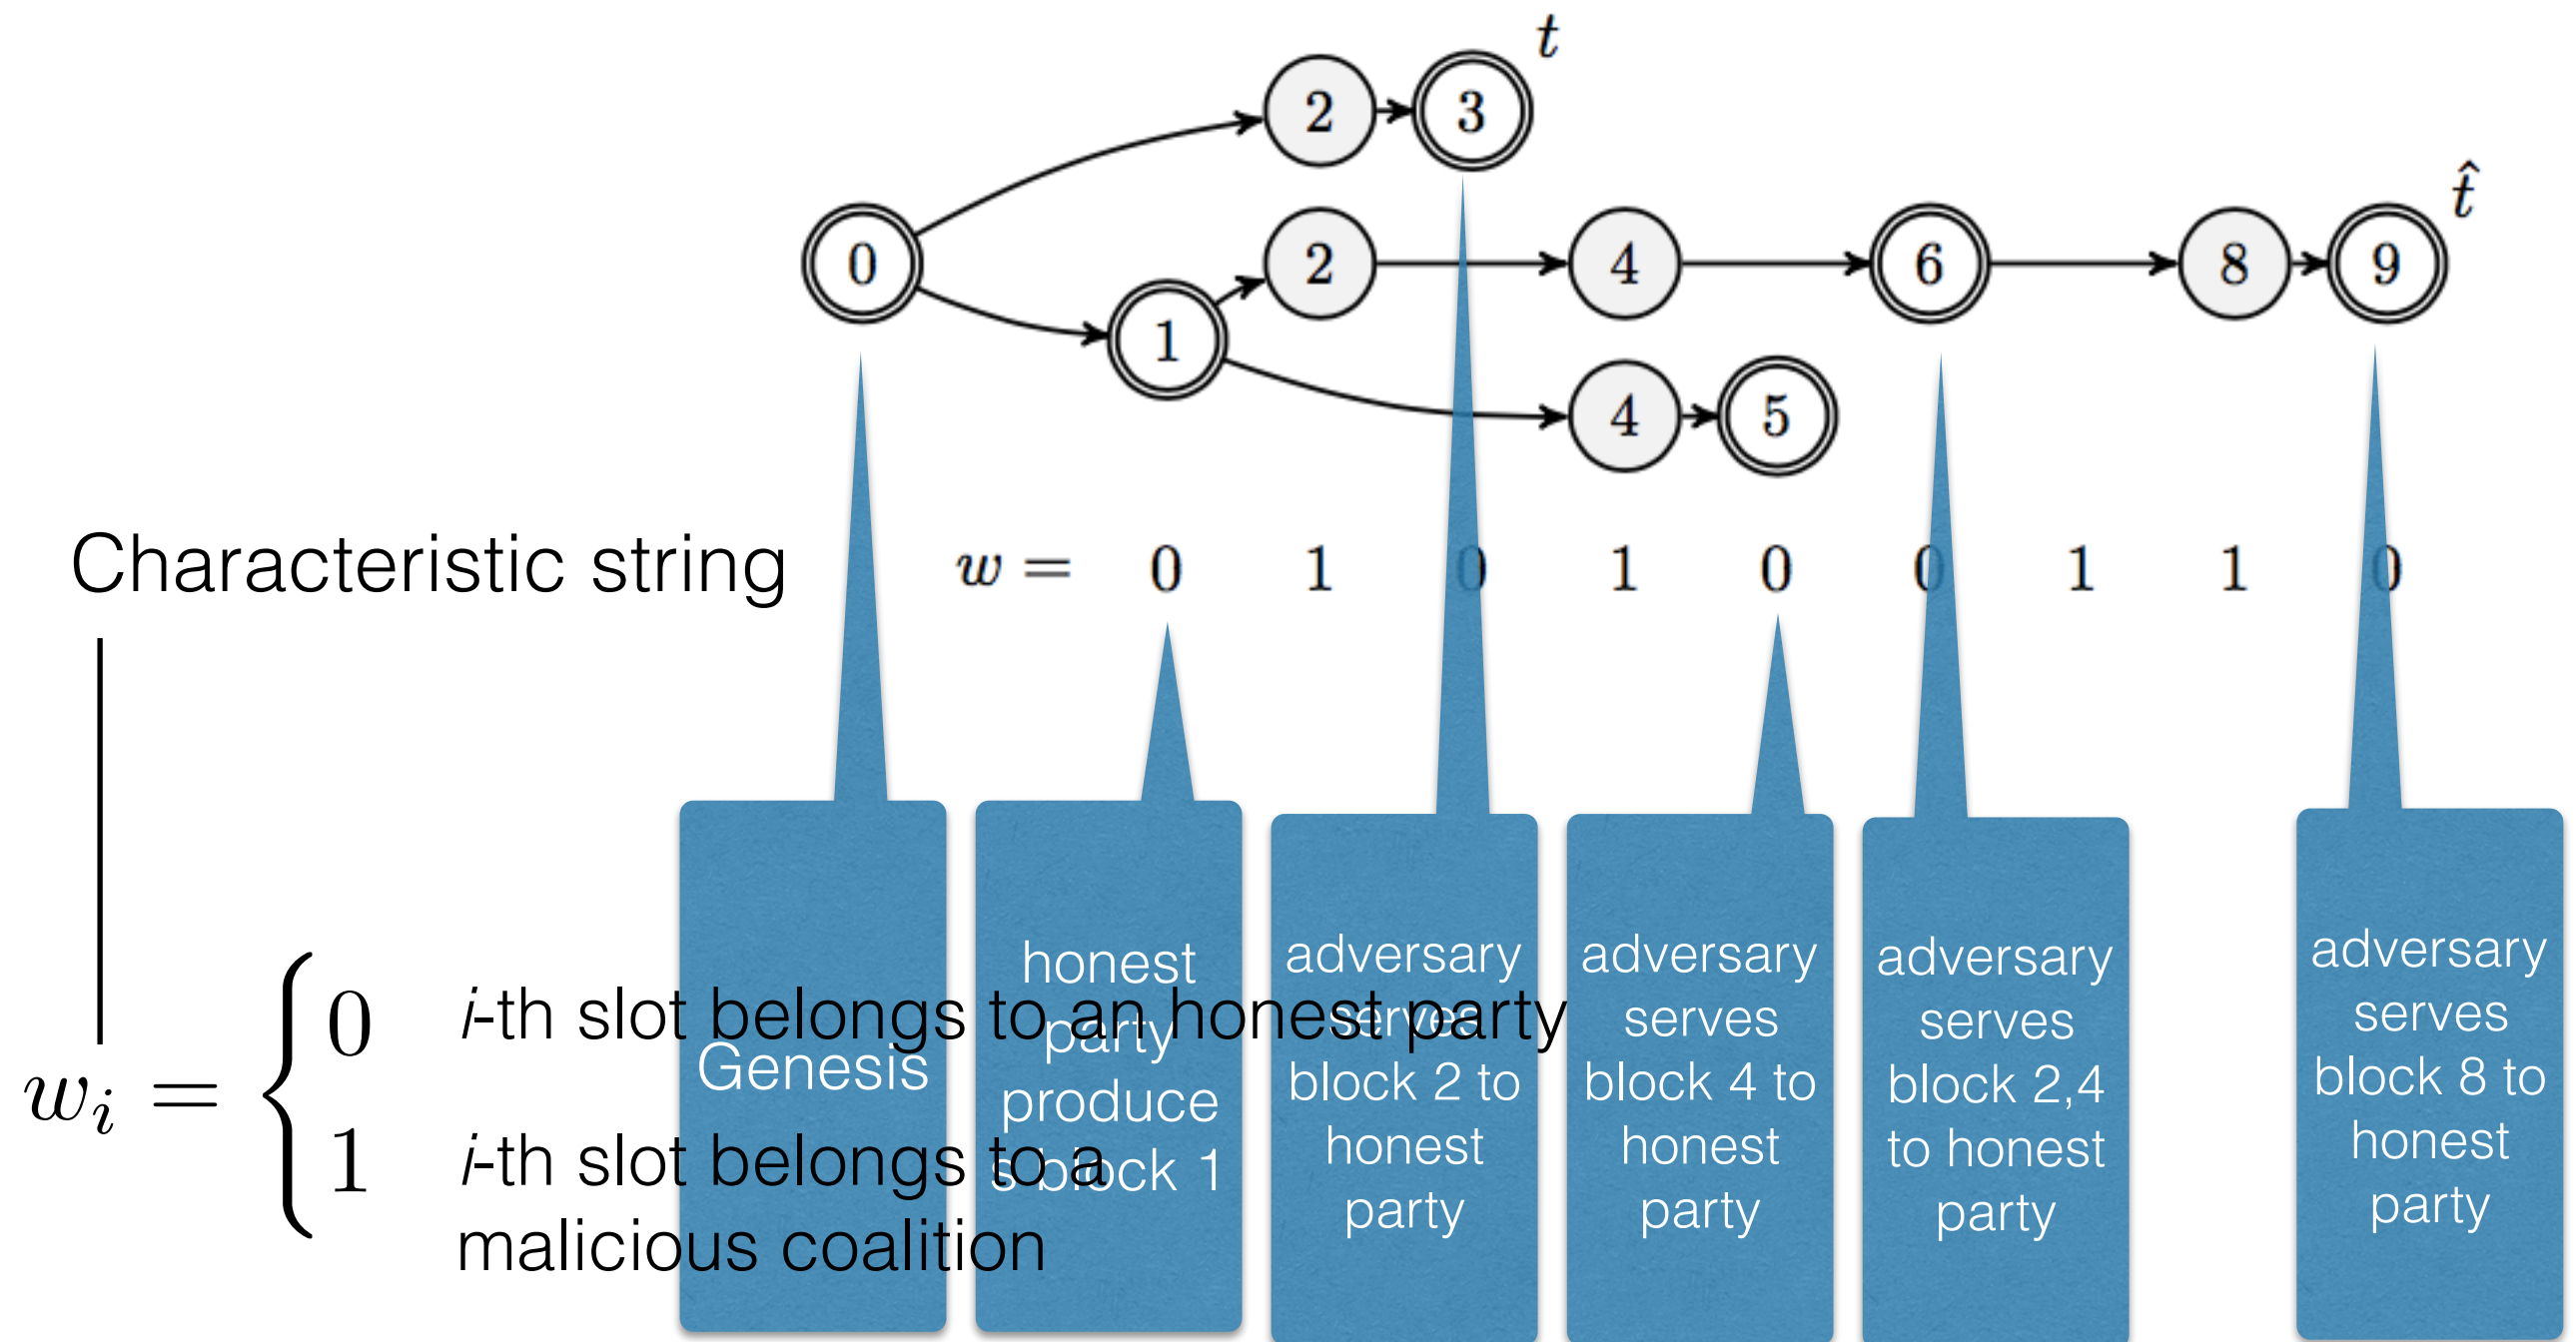
- Recall Proof of stake (PoS) protocols generalize the bitcoin blockchain protocol. Substitute PoW for PoS.
- Resulting protocol resembles blockchain operation (as opposed to the BFT approach).
- Example: the Ouroboros protocol.

# Ouroboros-BFT





# Forks and Protocol Executions




# PoW vs. PoS Proof

- The adversary is at a **much better position** in this protocol execution compared to Bitcoin's PoW-based execution.
- it can **see** ahead of time how stakeholders are activated.
- it can **generate multiple** different blocks for the **same slot** at any time **without cost**.
- it can **wait** and act just before an honest party comes online.


# Forkable Strings

The characteristic strings the adversary prefers!

Not Forkable!



Forkable!



# the $t < n/3$ setting

- Consider any string with Hamming weight less than  $n/3$ .
- **Lemma.** A slot controlled by the attacker can be used in at most two different locations to contribute blocks in the set of all admissible blockchains (those that correspond to protocol executions).
- It follows if  $t < n/3$ , there is not enough slots possessed by the adversary to match the blocks of the honest participants.

# Permissioned Ledgers with Dynamic participation

- Participation is still restricted: Producing transactions and/or blocks can only be performed only after being authorized.
- But dynamic: set of nodes implementing the protocol is evolving over time.
- Consider that each participant has some stake in the system

# Design Approach

- Iterate a basic recursive step:
  - Form a committee that will run a fixed term of ledger updates.
  - Update committee based on the outcome of the inputs posted on the ledger.

# Forming a committee, I

- The elected committee to perform the permissioned protocol execution should respect the adversarial bound under which the underlying protocol execution operates.
- Example of the simplest possible case: the stakeholder distribution is flat; a bound of  $t < n/3$  is required. The committee selected is the set of all stakeholders.
- A bound of  $< 1/3$  in terms of relative adversarial stake directly translates to a  $t < n/3$  bound in the selected committee.

# Forming a committee, II

- Consider a biased stakeholder distribution.
- Where the  $i$ -th stakeholder has stake  $s_i$  with total stake is  $S \in \mathbb{N}$
- Execute the permissioned protocol with the  $i$ -th stakeholder maintaining  $s_i$  identities out of a total space of identities that has cardinality equal to  $S$ .



# Forming a committee, III

- What if the size of the total stake  $S$  is prohibitively large?
- Random sampling may be used to select a subset of stakeholders to form the committee.
- Perform the sampling following “weighing by stake”: the probability that the  $i$ -th stakeholder is selected should be  $s_i / S$ .
- Randomness should be refreshed in each permissioned execution.

# Stake Shift

- The stake shift signifies the statistical distance between the random weighed by stake sampling of a stakeholder between two stakeholder distributions.
- E.g., if a stakeholder distribution has 3 members sharing equal stake and one of them is replaced, the stake shift is  $1/3$ .
- The maximum stake shift is 1 (the set stakeholders have completely changed).

# Refreshing Randomness

- Critical for security :
  - Possible via techniques from secure multiparty computation. Coin flipping protocols.
  - Or derivation from some aspect of the protocol execution (e.g., *hash* the previous block of transactions and use the output as coins to seed the sampling).
- Beware: **grinding attacks**; the attacker tries to impose a protocol history that favors adversarial parties.
  - *Hashing* a large sequence of previous values can be seen to minimise the impact of such attacks.

# New Parties Joining

- When a new party joins the system what information should be made available?
- **necessary**: the initial stakeholder distribution.
- **sufficient**: the most recent block for which consensus has been reached (checkpointing).

# Genesis Block Joining

- The setting when new parties joining have access only to the genesis block.
- One has to deal with a “**long range**” attack:
  - Attacker builds a LOG that is acceptable starting from the genesis block.
- New parties joining have the task to choose which is the correct version of the LOG.

# An example of a long range attack

- Assume the adversary can break with probability  $\epsilon$  the randomness generation and choose the outcome.
- Then it is impossible to protect against a long range attack: over an expected period of  $\epsilon^{-1}$  permissioned executions, the attacker will get the opportunity to transfer control to a “**sybil**” committee.
- (A sybil attack is when an attacker creates multiple fake identities that make the attacker appear as multiple parties in the system).

# Mitigating Long Range Attacks

- ensure good randomness generation (if random sampling is used).
- Identify characteristics in the LOG structure that signify that the LOG is not the correct one. For instance here are some ideas that have been proposed:
  - blocks are signed by too few of the stakeholders.
  - the sequence of blocks is too sparse in the time-domain in the PoS-based approach.

# End of lecture 08

- Next lecture:
  - Pitfalls and security vulnerabilities in smart contracts.